

Banner appropriate to article type will appear here in typeset article

Strain-coupled one-dimensional turbulence: a single-line closure for rapid distortion, its measured limits, and the consequence for leading-edge noise

Sparsh Sharma¹ and Alan R. Kerstein²

¹German Aerospace Center (DLR), 38108 Braunschweig, Germany

²Consultant, Danville, CA 94526, USA

Corresponding author: Sparsh Sharma, sparsh.sharma@dlr.de

(Received xx; revised xx; accepted xx)

XXXX

Key words: one-dimensional turbulence; leading-edge noise; turbulence modelling

1. Introduction

The distortion of turbulence by a mean velocity gradient is among the most basic problems of turbulence dynamics. When a turbulent flow is strained—accelerated through a contraction, deflected around a bluff body, compressed towards a stagnation region—the mean deformation reorganises the fluctuating field: it amplifies some velocity components and suppresses others, redistributes energy among them through the fluctuating pressure, and rescales the spectral content of the motion. When the strain accumulates in a time short compared with the eddy turnover time, the nonlinear interactions cannot keep pace and the response is governed, to leading order, by the linear dynamics of rapid distortion, first analysed by Batchelor & Proudman (1954) and Townsend (1976), developed for flow around obstacles by Hunt (1973) and Hunt & Graham (1978), and given its modern systematic form by Hunt & Carruthers (1990) and Sagaut & Cambon (2018). The turbulence that emerges from a strained region differs from the turbulence that entered it—in intensity, anisotropy and spectral shape—and that difference is often the quantity of practical interest, from grid turbulence in wind-tunnel contractions (Mills & Corrsin 1959; Tucker & Reynolds 1968; Lee & Reynolds 1985) to the distortion of oncoming turbulence approaching a lifting surface (Hunt & Graham 1978; Goldstein & Atassi 1976).

Capturing this distortion at a cost compatible with parametric study is difficult. Scale-resolving simulation represents it directly but precludes broad sweeps over intensity, length-scale ratio and Reynolds number; linear rapid-distortion theory (RDT) is exact in the rapid limit but neglects, by construction, the inter-scale transfer and the slow return to isotropy that act over any finite strain duration (Sagaut & Cambon 2018); second-moment closures

restore the nonlinear relaxation but discard the spectral information that the distortion redistributes (Launder *et al.* 1975; Pope 2000). Between these sits a gap: a reduced-order model that *evolves* strained turbulence with full scale resolution, retaining the rapid distortion and the nonlinear relaxation together, at negligible cost.

A concrete and demanding instance, and the application motivating this work, is turbulence–aerofoil interaction noise. The prevailing prediction framework (Amiet 1975; Paterson & Amiet 1977; Roger & Moreau 2010) obtains the radiated spectrum as a linear functional of the upwash velocity spectrum of the incoming turbulence—but the spectrum it consumes is prescribed, not computed, usually as an isotropic von Kármán form (von Kármán 1948), whereas the turbulence actually reaching the leading edge is strained turbulence. The importance of this pre-impact distortion for the prediction is by now well established: measurements and scale-resolved simulations show the near-edge spectrum systematically altered (dos Santos *et al.* 2023; Ribeiro *et al.* 2023; Piccolo *et al.* 2024), and corrected predictions have been obtained by prescribing RDT-transformed inflow spectra (de Santana *et al.* 2016; dos Santos *et al.* 2023) or by sampling the spectrum close to the edge (Piccolo *et al.* 2024). What these corrections share is that the distorted spectrum is imposed—from linear theory, or from a measurement that a prediction tool does not have. Synthetic-turbulence generators (Kim & Haeri 2015) share the limitation by construction: they reproduce a prescribed spectrum rather than evolving one. A model that predicts the distorted spectrum, including the nonlinear relaxation that linear transforms omit, would supply an element that is at present missing—here and wherever distorted turbulence feeds a downstream prediction, from contraction design to the turbulence ingested by rotors. The subject of this paper is that distortion problem; the acoustic stage enters only as the established linear functional that consumes the model’s output.

Map-based stochastic one-dimensional models offer the route. The one-dimensional turbulence (ODT) model (Kerstein 1999; Kerstein *et al.* 2001), built on the triplet map of the linear-eddy model (Kerstein 1991), evolves a three-component velocity field along a notional line, resolving the full range of length scales while representing turbulent advection as a stochastic sequence of mapping events interleaved with a continuous, one-dimensional specialisation of viscous transport; at each event a kernel mechanism redistributes energy among components, emulating pressure scrambling (Wunsch & Kerstein 2001; Kerstein *et al.* 2001).

Generalised to variable density (Ashurst & Kerstein 2005), adaptive meshes (Lignell *et al.* 2013; Stephens & Lignell 2021) and cylindrical geometry (Lignell *et al.* 2018), embedded as a subgrid closure (Schmidt *et al.* 2010), and coupled to acoustic analogies for jet noise (Sharma *et al.* 2022; Medina Méndez *et al.* 2023), ODT is an accurate and inexpensive generator of scale-resolved turbulence statistics. But standard ODT cannot represent distortion by a mean strain: its pressure treatment models only the slow, turbulence-driven return to isotropy, not the rapid part driven directly by the mean strain, and the apparent remedy—importing the RDT amplitude equation, whose pressure response is the solenoidal projection in three-dimensional wavenumber space—is unavailable on a line that resolves only the wavenumber along its own coordinate. Reconciling the inherently three-dimensional, non-local pressure physics of rapid distortion with the one-dimensional, local structure of ODT is the central modelling problem addressed here.

The paper reports six results:

- A strain-coupled formulation of ODT (SC-ODT) (§ 2): production enters as an exact continuous forcing, the rapid pressure–strain as an energy-conserving redistribution operator built from an explicit spectral closure, and the wavevector kinematics as a dilatation of the domain. The construction preserves the model’s structural invariants

and, under the stated closure, reproduces the rapid-distortion evolution of the component energies with no adjustable constant (§ 3).

- Validation against the strained-turbulence experiment of [Chen et al. \(2006\)](#) (CMK), summarised in § 4.5 with details in the supplementary material (§ S1).
- The linear reference appropriate to the model’s observables: exact RDT projected onto a line is broadband rather than a rigid spectral translation, so the reference against which any line model must be judged already differs from the model’s own kinematics at the linear level (§ 4.2).
- An independent spectral benchmark consisting of a strained-box direct numerical simulation together with exact-RDT companion calculations, reduced to the same one-dimensional observables. A three-way comparison localises the model’s remaining deficiency in its wavenumber-uniform rapid kinematics and supplies a distance-from-linear-theory envelope as a function of strain rapidity Sk_t/ε (§ 4.6).
- A treatment of what a single line can determine about the rapid pressure–strain term (a non-local three-dimensional object). The obstruction is stated exactly, sharp bounds are obtained from line data alone, and the axisymmetry hypothesis is measured rather than assumed: it is exact for axisymmetric distortion and carries a quantified error under plane strain (§ 5).
- The measured limits of the model and a mechanism that extends them (§ 6). A transmission diagnostic shows the strain-induced anisotropy crossing two triplet-map generations (a factor of nine in scale) undiminished; a controlled family of interventions rules out acceptance gating and map reweighting as remedies. Hierarchical application of the energy-redistributing kernel to eddy sub-scales supplies scale-local relaxation at the scales of the imbalance. The mechanism’s deep-strain fade is traced to its event-locked timing and is removed by an independent, size-capped relaxation clock, leaving a single measured allocation bound.
- The consequence for the motivating application (§ 7): the computed distorted spectrum is supplied to Amiet’s leading-edge-noise theory in place of the frozen inflow of standard practice, both as a ratio from which aerofoil geometry and observer cancel and as an absolute prediction at the reference operating point. The calculation quantifies the 5–20 dB term omitted by prescribed-spectrum practice; the assembled chain is compared with four published experiments and lies within 1–2 dB rms where the facilities are cleanest.

Throughout we show where the model agrees with reference data and, where it does not, by how much and for what reason.

The paper falls into three parts. The strain-coupled formulation of the first part modifies the model’s *continuous* advancement—the production and rapid pressure–strain of (2.8)—and leaves the discrete eddy events unchanged; the mechanism of the second part in turn modifies the *discrete* process, giving relaxation its own clock.

- The model and what a line can know: § 2 reviews standard ODT and derives the strain-coupled formulation; § 3 verifies it against exact rapid-distortion solutions and validates it at the moment level against the DNS of [Lee & Reynolds \(1985\)](#); § 4 constructs the linear reference, examines the emergent distorted spectrum and its operating point, and carries out the strained-box DNS benchmark; § 5 treats the single-line closure question.
- § 6 measures the model’s limits and establishes scale-local relaxation—timed by its own clock—as the mechanism that extends them.
- § 7 applies the computed distortion as a correction to the leading-edge-noise input.

Section 8 concludes. Supplementary material contains the validation against the strained-turbulence experiment of [Chen et al. \(2006\)](#) (CMK) and the complete set of intervention categories, labelled by category.variant (Appendix B).

2. Strain-coupled one-dimensional turbulence

The linear rapid-distortion equations, recalled below, contain two distinct ingredients: an evolution of the Fourier *amplitudes*, comprising production and the solenoidal pressure projection, and an evolution of the *wavevectors* that carries the kinematic straining of scales. The pressure part of the amplitude equation requires the full three-dimensional wavevector and is therefore not directly representable on a one-dimensional line; that incompatibility is the focus of this section. The distortion physics is split according to its structure and each part assigned to the mechanism that can carry it faithfully: the production, which is local in physical space, is imposed as a continuous forcing of the line velocity (§ 2.2); the rapid pressure–strain, which is not, is supplied as a continuous, energy-conserving redistribution operator constructed to reproduce the homogeneous rapid-distortion evolution of the component energies under an explicit spectral closure (§ 2.4); and the wavevector kinematics are recovered by a consistent straining of the domain (§ 2.5). The slow, turbulence-driven pressure–strain and the inertial cascade are captured by the discrete eddy events of the standard model, reviewed first.

Throughout, $A_{ij} = \partial U_i / \partial x_j$ denotes the imposed mean velocity gradient, taken uniform across the line at each instant; the medium is incompressible, $A_{ii} = 0$. The line carries the fluctuation field $u_i(y, t)$ about this mean, so that the total velocity is $U_i + u_i$ with $\overline{u_i} = 0$; $R_{ij} \equiv \overline{u_i u_j}$ is the Reynolds-stress tensor, the overline denoting the average over the statistically homogeneous directions (equivalently, the ensemble of independent realisations), and $k_t \equiv \frac{1}{2} R_{ii}$ the turbulent kinetic energy. In the homogeneous configurations used to establish and test the formulation the spatial line average and the ensemble average coincide.

2.1. The standard model

This subsection reviews the one-dimensional turbulence model of [Kerstein \(1999\)](#) in the vector-valued formulation of [Kerstein et al. \(2001\)](#), in the notation used throughout; nothing in it is original. The model evolves a three-component velocity field $u_i(y, t)$, $i = 1, 2, 3$, on a one-dimensional domain whose coordinate y is identified with the direction x_2 . The fields represent individual realisations; ensemble statistics are accumulated over many independent simulations. Two interleaved mechanisms advance the state (figure 1). Between events the fields obey ordinary one-dimensional diffusion,

$$(\partial_t - \nu \partial_y^2) u_i = 0, \quad (2.1)$$

with kinematic viscosity ν ; and a stochastic sequence of instantaneous *eddy events* represents the effect of individual turbulent overturns.

Because a continuum rotational motion cannot be represented on a line, an eddy event is an instantaneous, measure-preserving rearrangement of the profiles together with an additive velocity change,

$$u_i(y) \longrightarrow u_i(f(y)) + c_i K(y), \quad (2.2)$$

where f is the *triplet map* ([Kerstein 1991, 1999](#)): on an interval $[y_0, y_0 + l]$ of size l it compresses the profile to one third of its length, inserts three copies, reverses the central copy, and leaves the exterior unchanged. It is *measure preserving* (the non-local analogue

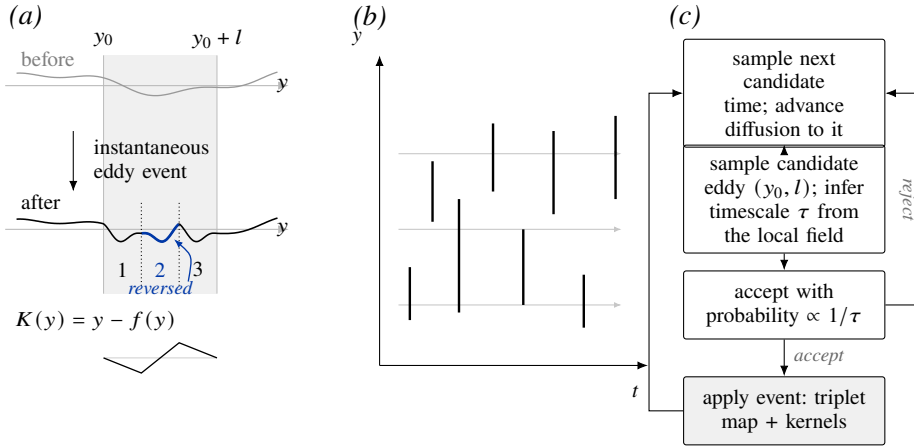


Figure 1. The model at a glance. (a) Instantaneous eddy event on $[y_0, y_0 + l]$: the triplet map compresses and triplicates the local field; the additive kernel $K(y)$ redistributes energy among components, both conserving momentum and total energy. (b) Eddy events (Poisson), the vertical bars marking the spatial range of each event, punctuating the continuous diffusion and mean-strain forcing (grey arrows). (c) Advancement cycle at a state-dependent rate; the $k^{-5/3}$ cascade emerges as an outcome of repeated sampling, consistent with inertial-range phenomenology.

of a divergence-free rearrangement), *continuous*, and *scale local* (it changes gradients by at most an order-unity factor) (Kerstein *et al.* 2001)—the three invariants whose preservation under an imposed mean strain is established in § 2.2. The additive term is built from the *kernel*

$$K(y) \equiv y - f(y), \quad \int K(y) dy = 0, \quad \int K^2(y) dy = \frac{4}{27} l^3, \quad (2.3)$$

non-zero only within the eddy interval; the first identity guarantees that the additive change conserves the momentum of each component, and the second fixes the energetics of the redistribution.

The kernel term is the model surrogate for the pressure-mediated redistribution of energy among velocity components (Wunsch & Kerstein 2001; Kerstein *et al.* 2001). Requiring that an event redistribute kinetic energy among components while conserving the total, with a rule invariant under permutation of the component indices, gives

$$\Delta E_i = \alpha \sum_j T_{ij} Q_j, \quad \mathbf{T} = \frac{1}{2} \begin{pmatrix} -2 & 1 & 1 \\ 1 & -2 & 1 \\ 1 & 1 & -2 \end{pmatrix}, \quad Q_i \equiv \frac{27}{8} \rho_0 l u_{i,K}^2, \quad (2.4)$$

where $u_{i,K} \equiv l^{-2} \int u_i(f(y)) K(y) dy$, Q_i is the available kinetic energy of component i within the eddy, ρ_0 the constant density, and $\alpha \in [0, 1]$ the parameter setting the fraction of extractable energy transferred towards isotropy. Solving the per-component energy balance for the amplitudes yields

$$c_i = \frac{27}{4l} \left[-u_{i,K} + \text{sgn}(u_{i,K}) \sqrt{u_{i,K}^2 + \alpha \sum_j T_{ij} u_{j,K}^2} \right], \quad (2.5)$$

real for $0 \leq \alpha \leq 1$. The transfer matrix $\mathbf{T}$ equalises the component energies: the construction is a model of the return-to-isotropy (*slow*) part of the pressure–strain interaction, a

correspondence made precise in § 2.3. The standard model contains no counterpart of the rapid, mean-strain-driven part.

The event sequence is a Poisson process whose rate density $\lambda(y_0, l; t)$ follows the standard eddy time-scale construction governed by the Q_i , with a rate constant, a viscous cutoff and a large-eddy suppression parameter as the only tunable constants; the inertial-range cascade and the $k^{-5/3}$ spectrum are *outcomes* of the sampling, not inputs (Kerstein 1999; Stephens & Lignell 2021). One diagnostic caveat of Kerstein *et al.* (2001) is respected throughout: the line velocity does not advect fluid, so diagonal component spectra and variances are read from the fields while shear stresses are read from eddy-induced transport.

2.2. The strain-coupled line equation and its structural admissibility

Standard ODT carries no mean strain: between events the field obeys the pure diffusion (2.1), and the eddy events supply only the turbulence-driven cascade and return to isotropy. The strain-coupled formulation generalises the *continuous* advancement to carry the distortion by an imposed mean velocity gradient A_{ij} , which is prescribed—the input to the problem—rather than evolved. Two irrotational homogeneous strains are studied, both admitting exact rapid-distortion solutions and isolating the physics of interest. *Plane strain*,

$$\mathbf{A} = \text{diag}(a, -a, 0), \quad a > 0, \quad (2.6)$$

is the strain of a two-dimensional stagnation region: with x_1 tangential to the surface, x_2 normal to it and x_3 spanwise, the component compressed by the strain is the wall-normal (upwash) component u_2 , whose amplification is the pre-impact distortion relevant to leading-edge noise. *Axisymmetric strain*,

$$\mathbf{A} = \text{diag}(-\tfrac{1}{2}\gamma, -\tfrac{1}{2}\gamma, \gamma), \quad \gamma > 0, \quad (2.7)$$

is the strain of a contraction or a three-dimensional stagnation flow and provides an independent test in which two components are amplified and one suppressed. For both, $\mathbf{A}$ is symmetric, so the principal-axis alignment assumed in § 2.5 holds, and the line is taken along x_2 . The modelling task is to add to the continuous advancement the two processes through which A_{ij} acts—an exact production and a modelled rapid pressure–strain—while leaving the eddy events unchanged.

The production term of the fluctuation momentum equation, $-A_{ij}u_j$, is a pointwise linear combination of the local velocity components and is therefore exactly representable on the line. Between eddy events the molecular evolution is replaced by

$$\partial_t u_i = \nu \partial_y^2 u_i - A_{ij} u_j + B_{ij} u_j, \quad (2.8)$$

the line form of the fluctuation momentum equation for a spatially uniform mean gradient: molecular diffusion as in standard ODT, the exact production $-A_{ij}u_j$ (whose second moment is the production $\mathcal{P}_{ij}$ of (2.10)), and, in place of the fluctuating pressure gradient, the modelled rapid pressure–strain operator $B_{ij}u_j$ constructed in § 2.4. The eddy events (2.2)–(2.5) are retained without modification, and the advancement is split between the continuous phase (2.8) and the discrete event sequence.

The forcing preserves the structural properties on which the model rests; we record the argument as a consistency check. Measure preservation, scale locality and the kernel identities (2.3) are geometric properties of f and K on a fixed domain, and the operator splitting applies the discrete event exactly as in the standard model, so they hold verbatim. For a mean gradient uniform along the line the operator $\mathcal{A}_{ij} \equiv -A_{ij} + B_{ij}$ acts pointwise, so over a non-diffusive sub-step the profile is advanced by a y -independent matrix, $u_i(y, t) = [\exp \int_{t_0}^t \mathcal{A} dt']_{ij} u_j(y, t_0)$ —continuous profiles therefore map to continuous profiles while

the molecular term is smoothing; and the forcing displaces no fluid along y , leaving the domain measure untouched (the deliberate domain straining is the separate operation of § 2.5).

The energetic role of the forcing is exact. At fixed domain length the production term evolves the Reynolds stress as

$$\left. \frac{dR_{ij}}{dt} \right|_{\text{prod}} = -A_{ik}R_{kj} - A_{jk}R_{ki} \equiv \mathcal{P}_{ij}, \quad (2.9)$$

exactly the Reynolds-stress production tensor of (2.10): the term carries no model assumption. The kinetic-energy budget of the formulation thus separates into an exact production by the mean strain, energy-conserving redistribution (the kernels), and dissipation by molecular diffusion—the structure of the exact transport equation. The forcing adds the one term the standard model lacks without disturbing the terms it already represents.

2.3. Rapid and slow pressure–strain, and linear rapid distortion

The production in (2.8) is exact; the rapid operator B_{ij} is not. Constructing it requires the second-moment and rapid-distortion physics it must reproduce, set out here. The Reynolds-stress transport equation for a fluctuating field about a mean with constant gradient A_{ij} reads, in homogeneous turbulence,

$$\frac{dR_{ij}}{dt} = \mathcal{P}_{ij} + \Pi_{ij} - \varepsilon_{ij}, \quad \mathcal{P}_{ij} = -A_{ik}R_{kj} - A_{jk}R_{ki}, \quad \Pi_{ij} = \left\langle \frac{p'}{\rho} (\partial_j u'_i + \partial_i u'_j) \right\rangle, \quad (2.10)$$

with ε_{ij} the dissipation tensor (Pope 2000). The fluctuating pressure satisfies a Poisson equation whose source splits into a part linear in the mean gradient and a part quadratic in the fluctuations (Pope 2000),

$$\frac{1}{\rho} \nabla^2 p' = \underbrace{-2 A_{lm} \partial_l u'_m}_{\text{rapid}} - \underbrace{\partial_l \partial_m (u'_l u'_m - \langle u'_l u'_m \rangle)}_{\text{slow}}, \quad (2.11)$$

and the pressure–strain correlation inherits the decomposition $\Pi_{ij} = \Pi_{ij}^{(r)} + \Pi_{ij}^{(s)}$ into a *rapid* part driven by the mean strain and a *slow* part driven by the turbulence alone. The slow part relaxes the Reynolds stress towards isotropy—the mechanism that the kernel (2.4)–(2.5) reproduces in model form. The rapid part is linear in A_{lm} and is determined by the velocity spectrum tensor $\Phi_{ij}(\kappa)$; the standard model represents neither $\mathcal{P}_{ij}$ nor $\Pi_{ij}^{(r)}$.

When the strain outruns the eddy turnover—the rapid limit—the response is computed exactly by linear theory (Batchelor & Proudman 1954; Hunt 1973; Hunt & Carruthers 1990). Writing the fluctuation as Fourier modes with wavevectors convected by the mean flow, each mode of time-dependent wavevector $\kappa(t)$ and amplitude $\hat{u}_i$ evolves, in the inviscid rapid limit, according to

$$\frac{d\kappa_i}{dt} = -A_{ji} \kappa_j, \quad \frac{d\hat{u}_i}{dt} = \left(\frac{2 \kappa_i \kappa_m}{\kappa^2} - \delta_{im} \right) A_{mn} \hat{u}_n. \quad (2.12)$$

The bracketed operator combines the production $-A_{im}\hat{u}_m$ with the rapid pressure response, the solenoidal projection $P_{ij} = \delta_{ij} - \kappa_i \kappa_j / \kappa^2$ applied to the strained field; the factor of two is fixed by continuity, since only that coefficient makes $d(\kappa_i \hat{u}_i)/dt$ vanish under (2.12). Two features are central to what follows. First, the pressure response is *non-local* and intrinsically *three-dimensional*: it requires all of $\kappa_1, \kappa_2, \kappa_3$, whereas a single line resolves only the wavenumber along its own coordinate, so (2.12) cannot be evaluated on the line

as written. Second, the rapid/slow decomposition furnishes the organising principle of the formulation: the slow part is already carried by the eddy kernel and the production by (2.8); what remains is to construct the rapid redistribution B_{ij} consistent with (2.12) from quantities available on the line—the task of § 2.4.

2.4. The rapid redistribution operator

Were the continuous evolution to consist of production and diffusion only ($B_{ij} = 0$), the line would reproduce $dR_{ij}/dt = \mathcal{P}_{ij}$, whereas the exact inviscid rapid-distortion balance is $dR_{ij}/dt = \mathcal{P}_{ij} + \Pi_{ij}^{(r)}$. The omission is not small: for turbulence isotropic at the onset of straining the classical result recovered below gives $\Pi_{ij}^{(r)} = -\frac{3}{5}\mathcal{P}_{ij}$, so production acting alone overstates the initial distortion rate by a factor of $\frac{5}{2}$; under sustained strain it also lets a component energy grow without the bound that incompressibility imposes. Representing $\Pi_{ij}^{(r)}$ is necessary.

2.4.1. The exact rapid pressure–strain and the closure problem

The exact rapid term follows from the amplitude equation. Writing (2.12) as $d\hat{u}_i/dt = M_{in}\hat{u}_n$ with $M_{in} = -A_{in} + 2(\kappa_i\kappa_m/\kappa^2)A_{mn}$, and noting that the incompressible wavevector dynamics preserve the measure $d\mathbf{\kappa}$ (their phase-space divergence is $-A_{ii} = 0$), the spectrum tensor $\Phi_{ij} = \langle \hat{u}_i^* \hat{u}_j \rangle$ yields

$$\frac{dR_{ij}}{dt} = \int (M_{in}\Phi_{nj} + M_{jn}\Phi_{in}) d\mathbf{\kappa} = \mathcal{P}_{ij} + \Pi_{ij}^{(r)}, \quad (2.13)$$

in which the production separates from the pressure contribution,

$$\Pi_{ij}^{(r)} = 2A_{mn}(\mathcal{R}_{ijmn} + \mathcal{R}_{jimn}), \quad \mathcal{R}_{ijmn} \equiv \int \frac{\kappa_i\kappa_m}{\kappa^2} \Phi_{nj}(\mathbf{\kappa}) d\mathbf{\kappa}. \quad (2.14)$$

The tensor is traceless—by solenoidality $\kappa_n\Phi_{nj} = 0$, the trace integrand contains $\kappa_i\Phi_{ni} = 0$ —so the rapid term redistributes energy among components without changing the total, in keeping with the energy-conserving character of the kernel mechanism. Equation (2.14) also exposes the central difficulty: the rapid term is a functional of the spectrum tensor through the orientation-weighted integral $\mathcal{R}_{ijmn}$, not of the single-point Reynolds stress—the classical closure problem of the rapid pressure–strain, common to every single-point model (Crow 1968; Launder *et al.* 1975; Pope 2000) and intrinsic to the physics. A line carries R_{ij} and its own one-dimensional spectra, but not the distribution of energy over wavevector orientation that $\mathcal{R}_{ijmn}$ requires; a spectral-tensor closure is therefore unavoidable, and we make it explicit. (What a line’s data *do* determine about this integral, and how sharply, is the analytical question taken up in § 5.)

2.4.2. Isotropy-based closure

The minimal closure takes the spectrum tensor instantaneously isotropic, $\Phi_{nj} = \frac{E(\kappa)}{4\pi\kappa^2}(\delta_{nj} - \kappa_n\kappa_j/\kappa^2)$ with $\int E d\mathbf{\kappa} = k_t$. With the angular averages $\overline{e_ie_m} = \frac{1}{3}\delta_{im}$ and $\overline{e_ie_me_ne_j} = \frac{1}{15}(\delta_{im}\delta_{nj} + \delta_{in}\delta_{mj} + \delta_{ij}\delta_{mn})$ for $e_i = \kappa_i/\kappa$,

$$\mathcal{R}_{ijmn}^{\text{iso}} = k_t \left[\frac{4}{15}\delta_{im}\delta_{nj} - \frac{1}{15}(\delta_{in}\delta_{mj} + \delta_{ij}\delta_{mn}) \right], \quad (2.15)$$

and substitution into (2.14) gives, with $S_{ij} = \frac{1}{2}(A_{ij} + A_{ji})$ and $A_{nn} = 0$,

$$\Pi_{ij}^{(r),\text{iso}} = \frac{4}{5}k_t S_{ij} = -\frac{3}{5}\mathcal{P}_{ij} \quad (\text{isotropic } R_{ij}), \quad (2.16)$$

which is the exact response of initially isotropic turbulence to a weak rapid strain (Crow 1968) and the content of the isotropisation-of-production model with $C_2 = \frac{3}{5}$ (Pope 2000)—obtained here, with no adjustable constant, as the isotropic limit of (2.14). For anisotropic turbulence $\mathcal{R}_{ijmn}$ must be modelled as a tensor function of R_{ij} ; the general representation linear in the anisotropy is that of Launder *et al.* (1975), to which the construction reduces in form. We take the isotropic closure (2.15)—the isotropisation-of-production (IP) model—as the starting point, and assess the anisotropic correction in § 3.

2.4.3. Realisation on the line

Given the closed rapid term—symmetric and traceless—we require a pointwise operator B_{ij} whose action on the line reproduces it. Demanding that the operator evolve the Reynolds stress by exactly $\Pi_{ij}^{(r)}$ implies

$$(\mathbf{BR} + \mathbf{RB})_{ij} = \Pi_{ij}^{(r)}, \quad (2.17)$$

a continuous Lyapunov equation with, for positive-definite $\mathbf{R}$, a unique symmetric solution $\mathbf{B}$. The solution conserves kinetic energy automatically—contracting (2.17) gives $2B_{ij}R_{ij} = \Pi_{ii}^{(r)} = 0$, so the rate of working of the operator on the line vanishes. In the isotropic case $\mathbf{R} = \frac{2}{3}k_t\mathbf{I}$ it reduces to the closed form

$$B_{ij} = \frac{3}{5}S_{ij}, \quad \mathcal{A}_{ij} = -A_{ij} + \frac{3}{5}S_{ij}. \quad (2.18)$$

B_{ij} is evaluated from the running Reynolds stress at each sub-step and enters the continuous advancement (2.8) alongside the production. Phenomenologically it represents the part of the redistribution driven by the mean gradient, which proceeds continuously at the strain rate, in contrast to the slow return to isotropy that the eddy kernel supplies at the eddy rate. Strain-driven effects (production and rapid pressure–strain) are continuous and enter through $\mathcal{A}_{ij}$; turbulence-driven effects (the inertial cascade and the slow return to isotropy) remain the discrete events.

The sense in which the formulation is rapid-distortion consistent can now be stated precisely, again as a consistency check rather than a theorem. By (2.9) the production term contributes $\mathcal{P}_{ij}$ to dR_{ij}/dt and by construction (2.17) the operator contributes $\Pi_{ij}^{(r)}$; their sum is the exact balance (2.13) for the prescribed closure, with no adjustable constant, recovering (2.16) at the onset of distortion of isotropic turbulence. Two qualifications delimit this statement: the consistency holds at the level of the *component energies* and is *conditional on the spectral closure*; the isotropic closure is exact only at onset, and its accuracy for anisotropic states is that of the isotropisation-of-production approximation it reproduces. The *spectral* distribution of the distorted energy is not enforced by (2.8): the operator acts uniformly on every resolved wavenumber, and the distribution of energy across scales emerges from its interaction with diffusion, the cascade and the domain straining. It is therefore a genuine prediction, tested in §§ 3–4—where the uniform, wavenumber-blind action of B_{ij} is also identified as the model’s measured spectral deficiency (§ 4.6.1).

2.5. Wavevector kinematics by domain straining

The second ingredient of (2.12), the wavevector evolution, carries the kinematic compression and stretching of scales and is independent of the amplitude evolution: in the method-of-characteristics reading of linear theory, the wavevectors are the characteristics along which the amplitudes evolve. On the line the relevant wavenumber is κ_2 . We restrict attention to a line aligned with a principal axis of the mean strain, so that $A_{12} = A_{32} = 0$ and a material element of the line remains a coordinate line—as holds for the irrotational

mean flow of a stagnation region. Then $d\kappa_2/dt = -A_{22} \kappa_2$, and since a material element of length ℓ stretches as $\dot{\ell}/\ell = A_{22}$ while wavenumbers scale as $1/\ell$, a uniform dilatation of the domain reproduces the wavevector kinematics exactly:

$$\frac{1}{L} \frac{dL}{dt} = A_{22}, \quad \implies \quad \kappa_2(t) = \kappa_2(t_0) \exp\left(-\int_{t_0}^t A_{22} dt'\right), \quad (2.19)$$

with L the domain length, implemented within the Lagrangian adaptive-mesh advancement of the model (Lignell *et al.* 2013) by imposing the mean-flow stretching on the mesh.

This step is not redundant with the amplitude operator. The Reynolds stress is an intensive second moment and is invariant under the uniform dilatation: the dilatation relabels the scales that carry the energy without changing the component energies, while $\mathcal{A}_{ij}$ changes the component energies without relabelling scales. The two operations are the physical-space counterparts of the characteristic motion and the amplitude evolution in (2.12), and together they account for the mean straining without double counting. Two limitations are inherent to representing a three-dimensional kinematics on one axis. First, the transverse wavenumbers also strain but are not represented; their effect is subsumed in the spectral closure, which already supplies the orientation information the line lacks. Second, a rotational mean flow, or a line not aligned with a principal axis, tilts the material line out of the coordinate direction and lies outside the formulation.

2.6. Summary of the formulation

The strain-coupled model advances the line field over a sub-step by composing (a) a *continuous strain–diffusion phase on a dilating domain*—time advance $\partial_t u_i = \nu \partial_y^2 u_i + \mathcal{A}_{ij} u_j$ with $\mathcal{A}_{ij} = -A_{ij} + B_{ij}$, where the production is exact and B_{ij} is given by (2.17) from the running Reynolds stress and the spectral closure—the LRR closure of § 3, with the isotropic form (2.18) its onset limit—while dilating the domain by (2.19)—and (b) *discrete eddy events* sampled from the standard rate distribution: the triplet map and the slow kernel, unchanged. Step (a), the continuous advancement, carries the strain-driven physics (the exact production and the rapid pressure–strain); step (b), the discrete events, carries the turbulence-driven physics (the cascade and the slow return to isotropy). The strain coupling thus modifies only the continuous advancement and leaves the discrete events unchanged; the discrete process is itself modified later, by the relaxation mechanism of § 6. Standard ODT is recovered exactly when $A_{ij} = 0$.

The relative importance of the two steps is governed by the ratio of the mean strain rate $S \equiv (2S_{ij}S_{ij})^{1/2}$ to the turbulence relaxation rate,

$$\chi \equiv \frac{S k_t}{\varepsilon}. \quad (2.20)$$

For $\chi \gg 1$ the distortion accumulates faster than the turbulence can respond, the eddy events are negligible over the distortion time, and the formulation reduces to the rapid-distortion limit of step (a), exact for the component energies under the closure. For $\chi \ll 1$ the turbulence relaxes between increments of strain and the slow kernel dominates. The stagnation region of a leading edge is the intermediate regime $\chi = O(1)$, in which neither limit is uniformly valid and both mechanisms contribute; the fidelity of the formulation across this range, and of the emergent distorted spectrum, are assessed against rapid-distortion solutions and scale-resolving data in §§ 3–4.

3. Verification, and moment-level validation

We separate two questions. *Verification* asks whether the formulation and its implementation reproduce a known exact solution in the regime where one exists—the rapid-distortion limit, tested analytically and by the numerical solution below. *Validation* asks whether the model reproduces the behaviour of real strained turbulence outside that limit; the spectral validation, against the strained-box DNS with its exact-RDT companion, is the subject of § 4, and this section closes with the moment-level validation against the strained-turbulence DNS of Lee & Reynolds (1985). Every test in this section operates at the Reynolds-stress level, by design; no spectral claim rests on it, and none is made before the spectral benchmarks of § 4.

One point of identification should to be kept in view throughout this section and the next. The model under test is the strain-coupled formulation (SC-ODT); *standard* ODT is its $A_{ij} = 0$ limit (§ 2.6) and, having no mean-strain coupling—neither the production forcing, the rapid pressure–strain operator, nor the domain dilatation—produces none of the strained curves shown here. Every departure from isotropy in what follows is generated by the added strain coupling, not by the eddy dynamics of the standard model, which supplies only the cascade and the slow return to isotropy. Where the figures refer to “SC-ODT” they mean this formulation; “standard ODT” is its strain-off ($A_{ij} = 0$) limit, and the relaxation mechanism of § 6 (“clock-relaxed ODT”) does not enter until Part II.

3.1. The exact onset rate

The two canonical strains of § 2.2—plane strain (2.6) and axisymmetric strain (2.7)—admit exact rapid-distortion solutions and isolate the physics of interest.

For turbulence isotropic at the onset of straining the continuous evolution (2.8) gives, by (2.9) and (2.16),

$$\left. \frac{dR_{ij}}{dt} \right|_0 = \mathcal{P}_{ij} + \Pi_{ij}^{(r),\text{iso}} = \frac{2}{5} \mathcal{P}_{ij} = -\frac{8}{15} k_t S_{ij}, \quad (3.1)$$

the exact rapid-distortion rate for initially isotropic turbulence, with no adjustable constant. This is a closed-form prediction the model must satisfy identically. Evaluated for the two canonical strains it gives, under plane strain, onset rates $\mp \frac{8}{15} k_t a$ for the stretched and compressed components with the neutral component unchanged, and, under axisymmetric strain, $+\frac{4}{15} k_t \gamma$ for each lateral component against $-\frac{8}{15} k_t \gamma$ for the axis. For plane strain the predicted upwash growth rate $+\frac{8}{15} k_t a$ is to be compared with the $\frac{4}{3} k_t a$ that production alone would give—the factor- $\frac{5}{2}$ overstatement anticipated in § 2.4. A model that passes (3.1) for both strains has its production and rapid terms verified jointly; failure would indicate an error in the operator $\mathcal{A}_{ij}$ or its implementation, isolated from any closure or cascade effect.

3.2. Finite total strain, and the necessity of the rapid term

Beyond onset the turbulence becomes anisotropic and the rapid term (2.14) ceases to be a function of the Reynolds stress alone, so the verification becomes a test of the spectral closure rather than of the construction. We integrate the model’s second-moment (Reynolds-stress) equation along each canonical strain with two closures: the isotropisation-of-production closure (2.16) (IP, with $C_2 = 3/5$) and the quasi-isotropic model of Launder *et al.* (1975) (LRR-QI, abbreviated LRR),

$$\Pi_{ij}^{(r)} = C_2 k_t S_{ij} + C_3 k_t (b_{ik} S_{jk} + b_{jk} S_{ik} - \frac{2}{3} b_{mn} S_{mn} \delta_{ij}) + C_4 k_t (b_{ik} W_{jk} + b_{jk} W_{ik}), \quad (3.2)$$

with anisotropy $b_{ij} = R_{ij}/2k_t - \frac{1}{3}\delta_{ij}$, mean rotation $W_{ij} = \frac{1}{2}(A_{ij} - A_{ji})$ (zero for the irrotational strains considered), and $C_2 = 4/5$, $C_3 = 7/4$, $C_4 = 131/100$; the value $C_2 = 4/5$ is fixed by requiring that (3.2) reduce at isotropy to the exact $\frac{4}{5}k_t S_{ij}$ of (2.16), so LRR shares the constant-free onset slope. The benchmark is the exact rapid-distortion solution, obtained by integrating the spectral equations (2.12) over a Gauss–Legendre ensemble of initial wavevector directions on the unit sphere: for initially isotropic turbulence the modal dynamics depend only on direction, so the Reynolds stress is the spherical average of the evolved modal covariance. This integration is independent of the IP and LRR closures and reproduces the onset slope (3.1) to within 10^{-16} , verifying the rapid-distortion machinery itself.

Figure 2 reports the component-energy fractions against total strain. Three observations follow. First, at onset all curves for each component are tangent—the constant-free identity (3.1) is satisfied—and through $e \lesssim 1$ the closures are nearly indistinguishable from the exact solution. Second, as anisotropy accumulates the closures depart, LRR uniformly less than IP: for axisymmetric strain at $e = 4$ the lateral fraction is 0.498 (exact), 0.468 (LRR), 0.450 (IP), and the axial fraction 0.003, 0.063, 0.100. This gap is the spectral-closure error in isolation—the eddy events are inactive in the rapid limit—and it bounds the accuracy attainable from any single rapid operator. Third, for plane strain both closures miss a *qualitative* feature: in the exact solution the neutral spanwise component grows monotonically, overtakes the upwash near $e \approx 3.3$, and is the largest component by $e = 4$ ($\overline{u_3^2}/2k_t = 0.51$ against 0.47), whereas both closures predict it to decay. This redistribution into the neutral direction is governed by the orientation distribution of the spectral tensor and cannot be recovered from single-point quantities; it is the moment-level signature of precisely the transverse-wavenumber information a single line does not carry. Section 5.6 returns to it with the strained-box DNS in hand and localises the deficiency in the closure’s b -dependent term.

The upwash component, whose spectrum is the input to the leading-edge response, makes the case for the rapid term quantitative (figure 3). Production acting alone drives the upwash fraction from $1/3$ to 0.98 by $e = 4$ —almost the entire kinetic energy forced into one component, an essentially non-realisable state—and overstates the initial amplification rate by the factor $\frac{5}{2}$. The rapid term removes three fifths of this at onset and keeps the fraction physical: the exact solution peaks at 0.49 near $e \approx 2.4$ and then relaxes as energy passes to the spanwise component. Both closures reproduce the amplification itself—the effect that the frozen, isotropic input of classical leading-edge-noise prediction omits—while overpredicting it at large strain (LRR 0.60, IP 0.65 at $e = 4$, against the exact 0.47). The rapid pressure–strain is therefore not a refinement but a necessity, and the accuracy of the method for a given configuration is controlled by the total strain accumulated along the stagnation streamline.

3.3. Verification of the implementation

The preceding checks verify the *formulation*; the *implementation*—in which the operator is not prescribed but reconstructed at every substep (line Reynolds stress from the resolved field, closure, Lyapunov solve (2.17), pointwise application of $\mathcal{A}_{ij}$)—is verified in the supplementary material (§ S5, figure S5). Run in the rapid-distortion limit from a statistically isotropic resolved initial field, the solver lies on the LRR moment trajectory over $0 \leq e \leq 4$ to within 10^{-4} in each component fraction and 0.3% in the kinetic-energy amplification, so the per-substep chain is correct at every state of anisotropy; agreement with the *exact* solution would indicate a fault, since with events and diffusion suppressed the line carries no information beyond the single-point statistics the closure itself uses. The




Figure 2. Component energy fractions against total strain for initially isotropic turbulence under (a) plane strain and (b) axisymmetric strain, normalised to $S = 1$: exact rapid-distortion theory (wavevector-ensemble integration of (2.12)) against the SC-ODT second-moment equation with the LRR closure (3.2) and the IP closure (2.16). Line style denotes the source (solid exact RDT, dashed LRR, dotted IP) and marker the component ($\circ \overline{u_1^2}$, $\square \overline{u_2^2}$, $\triangle \overline{u_3^2}$). For axisymmetric strain the two lateral components coincide and only $\overline{u_1^2} = \overline{u_2^2}$ (lateral) and $\overline{u_3^2}$ (axial) are shown. All closures share the exact onset slope; LRR tracks the exact solution more closely at finite strain.




Figure 3. Upwash energy fraction under plane strain for initially isotropic turbulence: exact rapid-distortion theory, SC-ODT with the LRR and IP closures, and production alone ($B_{ij} = 0$).

test requires the resolved spectral initialisation and strain-resolved timestep bound recorded in Appendix B. For definiteness, *every strained computation in the remainder of the paper uses the LRR closure*, whose error against the exact solution is quantified in § 3.2 at every strain reached; the sensitivity of the spectral conclusions to this closure choice is not estimated but measured, through the exact-RDT comparisons of § 4 and the axisymmetry measurement of § 5, which carry the closure deficiency into the quoted uncertainty rather than leaving it implicit.

3.4. *Moment-level validation against the strained-turbulence DNS of Lee & Reynolds*

The DNS of Lee & Reynolds (1985) is retained as a benchmark for one specific reason: it was deliberately run at a turbulence Reynolds number low enough ($q^4/\nu\epsilon \lesssim 100$) to be fully resolvable on a single line, and it reports the Reynolds-stress anisotropy against total strain for several irrotational strain modes alongside the analytic rapid-distortion prediction—a graded, moment-level reference that the spectral benchmarks of § 4 do not provide. The same programme’s later rapid-limit analysis (Lee *et al.* 1986) supplies the analytic reference used in the CMK comparison; higher-Reynolds-number strained computations exist, but the spectral role they would serve is filled here by the dedicated strained-box DNS of § 4, leaving Lee & Reynolds’s low-Reynolds-number study as the appropriate moment-level benchmark for a fully resolved line. It is a Stanford report rather than a journal paper, but it remains the irrotational-strain benchmark of that programme’s own later work: the structure-based models of Reynolds & Kassinos (1995) and Kassinos *et al.* (2001) are validated against its plane-strain and axisymmetric runs (Kassinos & Reynolds 1998). More recent direct simulations of the same deformations at moderate Reynolds number provide independent, journal-published comparators, and are overlaid on the comparisons below: Zusi & Perot (2013) for plane strain (figure 4) and Zusi & Perot (2014) for axisymmetric contraction (figure 5), both digitised from their published anisotropy figures and shown over the strain range they reach ($c \leq 1.65$, where $c = \exp \int S dt$ is the deformation ratio, and $\ln c$ is the total strain e). A direct simulation of the CMK straining cycle (Gualtieri & Meneveau 2010), which reports the same variances and pressure–strain budget as that experiment against a Launder–Reece–Rodi closure, is discussed in the supplementary material. The comparison is based on the c dependence of the anisotropy b_{ij} as a function of the deformation ratio. computed from the single-point component variances along the line, ensemble-averaged over 1000 realisations; no spectral transform is involved. Two of their strain modes are run: plane strain and axisymmetric contraction, the line along a compressed direction in each case, to the total strains of the DNS ($c = 4$ and 3.32); the measured line compression agrees with $\exp(A_{22}t_f)$, where t_f is the total straining time, to better than 0.3%, confirming the kinematic fidelity of the dilatation. One parameter distinction matters for interpretation: Lee & Reynolds characterise each run by $S^* = Sq^2/\epsilon$, which equals 2χ identically since $q^2 = 2k_t$ —a correspondence of definitions, not a calibration—and find the axisymmetric anisotropy essentially independent of S^* —a clean, parameter-free target—whereas in plane strain the unstrained component b_{11} is strongly S^* -dependent, spanning $b_{11} \approx 0$ (slow) to $+0.12$ (rapid) in the DNS.

Under plane strain (figure 4) the model reproduces the two strained components quantitatively across the whole range: $b_{22} = +0.21$ at $c = 3.86$ against $\approx +0.20$ in the DNS, and $b_{33} = -0.20$ against ≈ -0.26 , capturing the trend and the bulk of the magnitude. These are exactly the components the DNS identifies as strain-rate-independent, so they constitute the robust part of the comparison. Under axisymmetric contraction the model reproduces the sign structure and the monotonic approach to a strained state while under-predicting the anisotropy magnitude by roughly 25% ($b_{11} = -0.21$ against ≈ -0.29 at $c = 3.16$), dominated by the spectral-closure gap isolated in § 3.2 and common to the whole second-moment family; the line adds only a small artefact, breaking the transverse symmetry with a residual split $b_{22} = +0.102$ against $b_{33} = +0.113$ that the symmetry-respecting mean $\frac{1}{2}(b_{22} + b_{33})$ removes. Figure 5 shows the comparison, with the independent DNS of Zusi & Perot (2014) tracking both the model and Lee & Reynolds over the strain range it reaches.

The one component the model does not reproduce, the plane-strain unstrained anisotropy b_{11} , is not a deficiency specific to the model, and figure 6 locates it exactly. Exact linear rapid distortion, evaluated by the wavevector-ensemble integration of § 3.2, gives b_{11} rising




Figure 4. Reynolds-stress anisotropy under plane strain: SC-ODT (lines; ensemble mean over 1000 realisations, standard error within the line width) against the DNS of Lee & Reynolds (1985) (open symbols, digitised from their figure 5.4). As noted in § 3.3, here and henceforth all strained computations use the LRR closure. Line style and marker denote the component (b_{11} solid/ $\circ$, b_{22} dashed/ $\square$, b_{33} dotted/ $\triangle$); open markers are Lee & Reynolds (1985), filled markers the independent DNS of Zusi & Perot (2013) (their highest rate, over the strained range $c \leq 1.65$); the hatched band is the DNS b_{11} range over slow-to-fast S^* . The strained components b_{22} and b_{33} are reproduced quantitatively, and the two DNS agree where they overlap; the unstrained b_{11} is strain-rate-dependent in the DNS, with SC-ODT at its slow-strain edge.




Figure 5. Reynolds-stress anisotropy under axisymmetric contraction: SC-ODT (lines; ensemble mean over 1000 realisations, standard-error band) against the DNS of Lee & Reynolds (1985) (open markers) and the independent DNS of Zusi & Perot (2014) (filled markers, their highest rate, over $c \leq 1.65$). Solid/ $\square$: the compressed axial component b_{11} ; dashed/ $\triangle$: the transverse component $\frac{1}{2}(b_{22} + b_{33})$. The model tracks the sign and monotonic growth of both, and the two DNS agree where they overlap.




Figure 6. The unstrained-direction anisotropy b_{11} under plane strain: SC-ODT (solid, 1000 realisations, grey standard-error band) against exact linear rapid distortion (long-dashed, wavevector-ensemble integration), the IP (dotted) and LRR (short-dashed) closures, and the DNS of [Lee & Reynolds \(1985\)](#) (open squares; hatched band = DNS range over slow-to-fast S^*). The fast-strain DNS tracks exact RDT, whose rapid pressure–strain feeds the neutral direction; the single-point closures—and SC-ODT, built on them—miss this orientation-carried redistribution, the deficiency measured and localised in § 5.6.

to $+0.114$ at $c = 4$: although $A_{1k} = 0$, the rapid pressure–strain feeds the neutral direction—the same orientation-carried redistribution as the spanwise growth in § 3.2. (A tempting kinematic argument, that $\langle u_1^2 \rangle$ is conserved because the component is unstrained, holds for production but not for the rapid term, and gives the wrong sign; an earlier version of this comparison used it.) The fast-strain DNS therefore *tracks exact RDT*, as [Lee & Reynolds](#) themselves observed, while both single-point closures give $b_{11} \approx 0$, slightly negative: the closure family misses the neutral-direction feed, which is carried by the orientation distribution that single-point models discard—the same deficiency measured spectrally, and localised in the closure’s b -dependent term, in § 5.6. SC-ODT coincides with the closure on which its operator is built, so its near-zero b_{11} is correct for its formulation and matches the DNS at slow strain rate, where the nonlinear scrambling erases the neutral-direction feed; at fast strain it inherits the closure family’s known limitation (§ 3.2) rather than any line-specific artefact. The components that govern the upwash spectrum fed to the leading-edge response are precisely those the model reproduces quantitatively. Finally, [Lee & Reynolds](#)’s own transport-budget analysis finds that in plane strain the intercomponent transfer is carried almost entirely by the rapid term—independent support for the rapid/slow decomposition on which § 5 rests.

4. The linear reference and the distorted spectrum

This section establishes what SC-ODT predicts for the strained spectrum and measures it against the correct linear reference and against DNS. It first separates the model’s own kinematics—a rigid spectral translation—from the prediction of linear theory for the line-projected spectrum, which is *not* rigid (§§ 4.1–4.2); it then examines the emergent distorted spectrum and fixes the physical operating point (§§ 4.3–4.4); and it closes with the




Figure 7. Compression without cascade (eddy events suppressed, inviscid). (a) Component line spectra $E_i(k_2)$ at total strains $e = 0, 1, 2, 3, 3.9$ (lighter to darker); the spectra translate to higher wavenumber without changing shape. (b) Spectral-centroid migration $\bar{\kappa}(e)/\bar{\kappa}(0)$ for each component (symbols) against the kinematic prediction $e^{-A_{22}e}$ (dashed). The centroid follows the prediction and the three components coincide, confirming purely geometric compression of the line.

benchmark that addresses the model-versus-itself comparison—a strained-box DNS with an exact rapid-distortion companion, reduced to the same one-dimensional observables (§ 4.6).

4.1. Compression without cascade: the model’s kinematic reference

With the eddy events suppressed and zero molecular viscosity, the mean strain is the only process acting on the resolved field, and its effect on the line is purely kinematic. The line length contracts as $\mathcal{L}(e) = \mathcal{L}_0 e^{A_{22}e}$, so every resolved wavenumber is rescaled as $k(e) = k(0) e^{-A_{22}e}$ —the wavevector kinematics it shares with rapid-distortion theory (Batchelor & Proudman 1954; Hunt & Carruthers 1990)—while the continuous redistribution operator of § 2.4 acts uniformly on every resolved wavenumber. The combination induces a rigid translation of the line spectrum towards higher wavenumber that transfer no energy between scales. We emphasise at the outset that this rigid translation is a property of the SC-ODT line—uniform dilatation composed with a wavenumber-uniform operator—and not of linear distortion theory itself; the correct linear reference for the projected spectrum is constructed in § 4.2 below. As a check on the implementation, however, the rigid translation is an exact, parameter-free target.

Figure 7 confirms it. The component spectra translate toward higher k without change of shape—the peak neither broadens nor develops an inertial range—and the spectral centroid follows $e^{-A_{22}e}$, reaching 7.03 at $e = 3.9$ against the exact 7.03, with the three components migrating identically; the domain length follows $e^{A_{22}e}$ to four digits. The single-point moments are unchanged from the no-dilatation case, since a uniform rescaling of position leaves the length-weighted velocity statistics invariant. Compression therefore acts entirely on the spectrum and reproduces the model’s kinematic prediction exactly. It is the kinematic reference (the dashed line in all centroid figures below) against which the full model is measured.

4.2. The linear reference for the projected spectrum: exact RDT

The rigid translation above must not be mistaken for the prediction of linear rapid-distortion theory for the quantity the line actually carries. In the full linear theory the amplitude of




Figure 8. Exact linear theory does not translate the projected spectra rigidly. (a) Shape ratio $R_{nn}(\kappa_2)$ of the exactly distorted, line-projected component spectra to a rigid translation under plane strain, at $e = 1$ (solid, filled markers) and $e = 0.5$ (dashed, open markers); markers distinguish the components, and the dotted line is the rigid translation $R_{nn} = 1$. Closed-form Cauchy solution for an isotropic initial spectrum, integrated over the transverse wavenumber plane. (b) Transverse splitting $\phi_{11}/\phi_{33} - 1$ versus wavenumber in exact linear theory (dashed) and in the strained-box DNS of § 4.6 (solid, squares), at $e = 1$ (heavy) and $e = 0.5$ (light).

each Fourier mode evolves at a rate that depends on the orientation of its three-dimensional wavevector (equation 2.12); the one-dimensional component spectra are projections—integrals over the transverse wavenumber plane—and there is no reason for a projection of orientation-dependent amplification to translate rigidly. Whether it does is a question of calculation, and the answer is that it does not. Using the closed-form Cauchy solution of the linear problem for an isotropic initial spectrum, integrated analytically over the transverse wavenumber plane and band-limited to the wavenumbers resolved by the DNS box of § 4.6, we define the shape ratio $R_{nn}(\kappa_2)$ as the projected component spectrum at strain e , normalised by its integral, divided by the rigidly translated initial spectrum normalised likewise; $R_{nn} \equiv 1$ at every κ_2 if and only if the translation is rigid. At $e = 1$ under plane strain the upwash ratio R_{22} falls monotonically from 0.86 at half the initial spectral centroid to 0.36 at ten times it: exact linear theory redistributes the projected upwash spectrum towards low wavenumber by a factor of 2.4 across the resolved range, because modes aligned with the line ($\kappa \rightarrow \kappa_2$) are amplified least. The transverse components are affected more weakly (R_{11} from 1.07 to 0.81; R_{33} within seven per cent of unity), so the projected anisotropy of linear theory is itself scale-dependent. Figure 8 shows the ratios. The exact-RDT companions of the DNS (§ 4.6), which share its initial fields, reproduce these analytic curves to within 0.05 in the transverse splitting over the resolved band; the analytic integral is used because the companions sample the largest scales with one Fourier line per band.

Two consequences run through the remainder of the paper. First, the comparison of the full model against the rigid translation (§ 4.3) measures the departure of the model from its own kinematics; the departure of the flow from linear theory must be measured against the exact projected reference, and the two differ already at the linear level. Second, the model’s rigid spectral kinematics is itself a structural approximation—linear theory deepens the projected anisotropy with wavenumber, which a wavenumber-uniform operator cannot represent—and § 4.6.1 quantifies the resulting deficiency against both the exact reference and DNS.

4.3. The distorted spectrum: strain on against strain off

The central model-internal comparison isolates what the strain does to the spectrum at fixed Reynolds number, holding the turbulence otherwise identical. Two ensembles of 128 realisations are run at $\nu = 10^{-5}$: SC-ODT (eddies and dilatation active) and a baseline identical in every respect except that the strain machinery is switched off, so the line does not compress—which is standard ODT, the $A_{ij} = 0$ limit of § 2.6. Because the two differ only in the presence of the mean strain, the difference between them is the distortion, cleanly separated from the cascade and dissipation the eddies produce on their own.

Figure 9 shows the result, and it departs from the kinematic reference in a specific and physically meaningful way. Consider the centroid first. The no-strain baseline, after the brief initialisation transient, relaxes monotonically to $\bar{\kappa}/\bar{\kappa}_0 < 1$: with no compression, the eddies merely redistribute energy and the centroid drifts below its initial value. SC-ODT is lifted clearly above this baseline—compression does push the spectrum towards higher wavenumber—yet it settles at $\bar{\kappa}/\bar{\kappa}_0 \approx 1.6$ –2.1, far below the rigid-translation line, which rises to 7.4 at $e = 4$. The reason is visible in the spectra: the SC-ODT spectrum is not a rigid translation of the baseline but a *broadband* distortion. It extends to substantially higher wavenumber than the baseline—energy is genuinely carried to small scales by compression and cascade—but the bulk of the energy remains at the energetic scales, and what reaches high wavenumber is dissipated there rather than accumulated.

Two distinct effects are mixed in this departure, and § 4.2 requires that they be separated. The first is the eddy cascade and dissipation, which the strain-on/strain-off construction isolates cleanly at the level of the model. The second is that exact linear theory itself, once projected onto the line, is already broadband (figure 8): part of what a rigid-translation reference attributes to nonlinearity is present in the linear physics. The strain-on/strain-off comparison therefore measures what the mean strain does *within the model*; how the model's distorted spectrum stands against exact linear theory and against DNS on the same footing is answered by the three-way comparison of § 4.6.1. What survives either way is the qualitative point that a frozen, rigidly strained upwash spectrum misrepresents the distorted inflow, which motivates supplying a distorted spectrum to Amiet's theory in place of it, in the spirit of the rapid-distortion-modified inflow spectra of de Santana *et al.* (2016) but without the frozen-shape assumption.

4.4. The physical operating point

The molecular viscosity is, in these units, the inverse integral-scale Reynolds number, and it is a physical property of the represented inflow, not a tunable parameter: a sweep over $\nu = 10^{-4}$ – 10^{-7} shows the cascade deepening monotonically, with the resolved cases settling below the linear-RDT centroid line while the lowest viscosities climb towards it only by piling energy into an under-resolved grid-scale shelf. The operating point is therefore the viscosity matching the target inflow, resolution-checked. For the reference configuration—a NACA0012 at $U_\infty = 20 \text{ m s}^{-1}$ in grid turbulence with intensity $u'/U_\infty \approx 6\%$ and integral scale $\mathcal{L} \approx 55 \text{ mm}$ at the leading-edge plane—the integral-scale Reynolds number is $Re_{\mathcal{L}} \approx 4.4 \times 10^3$, which maps through the initial-condition normalisation to $\nu \approx 3 \times 10^{-6}$, run on a refined grid (16000 cells) on which the dissipation rolloff is resolved with margin; the ensemble comprises 200 realisations.

Figure 10 shows the result. The component spectra fall smoothly across four decades with no grid-scale shelf, the upwash component E_2 amplified above the streamwise E_1 —the spectral signature of the plane strain—and no extended $k^{-5/3}$ range, as expected physically at this Reynolds number (Kolmogorov 1991; Pope 2000). The centroid settles after the initialisation transient and migrates to ≈ 2.4 by $e = 4$, far below the rigid-translation value




Figure 9. SC-ODT (strain on) against the matched strain-off baseline, standard ODT, at $\nu = 10^{-5}$ (ensemble means over 128 realisations, geometric $\pm\sigma$ bands). (a) Final-strain total spectrum $E(k_2)$: SC-ODT (strain on, solid) extends to higher wavenumber than the baseline (strain off, dashed) but remains broadband, not a rigid translation. (b) Centroid migration: the baseline relaxes below unity, SC-ODT is lifted above it by the compression, yet stays far below the rigid-translation line (grey dotted). The gap between the strain-on and strain-off curves is the distortion attributable to the mean strain at this Reynolds number.

of 7.4: the broadband, dissipative distortion of § 4.3 at the physical operating point. The regime is quantified with an estimator-free dissipation: because the eddy events conserve energy and the production is exact, molecular dissipation is the only sink, so it is fixed by the resolved energy budget

$$\varepsilon = \mathcal{P} - \frac{dk_t}{dt} = -A_{ij}R_{ij} - \frac{dk_t}{dt}, \quad (4.1)$$

every term available on the line. By this budget the strain-rate parameter settles at $\chi = Sk_t/\varepsilon \approx 1.2$ over the strained trajectory (ensemble of 200 realisations), corroborated within a few tens of per cent by the direct gradient estimate $\langle 2\nu(\partial u_i/\partial y)^2 \rangle \approx 1.1$: the operating point sits at $\chi = O(1)$, bracketed by the two strained-box DNS rapidities $Sk_t/\varepsilon = 0.8$ and 16 used in § 4.6.1, the regime the leading-edge stagnation flow occupies, where neither the rapid-distortion nor the near-equilibrium limit applies and where improvement over linear theory is meaningful. The spectral *peak* corroborates the regime reading independently of the dissipation estimator: the energy-containing wavenumber itself migrates far less than the rigid law, so the eddy relaxation holds back the energy-containing range, not merely the dissipation tail—linear theory misplaces the energy-containing wavenumber, not only the spectral shape at high κ . The operating-point ensemble completes in a few hundred core-hours, roughly two orders of magnitude below a strained DNS ensemble at comparable Reynolds number, which permits parametric sweeps over inflow condition.

4.5. Validation against the CMK strained-turbulence experiment

The strain-on/strain-off comparison is model-internal; an external benchmark at laboratory Reynolds number is provided by the CMK plane-strain experiment, whose straining phase (total deformation $D \approx 3$, $e \approx 2.20$ in the present notation) we reproduce. In the rapid-distortion limit (eddy events suppressed, inviscid) the model follows the analytic reference of Lee *et al.* (1986) that CMK use—component centroids on the rigid-translation law with deformation ratio 3.00 at $e = 2.20$, kinetic-energy amplification 1.72 against the exact 1.69, the residual attributed to the moment closure, not the implementation (§ 3.3). With




Figure 10. The physical operating point ($\nu \approx 3 \times 10^{-6}$, $Re_{\mathcal{L}} \approx 4.4 \times 10^3$, 16000 cells; ensemble over 200 realisations, geometric $\pm\sigma$). (a) Final-strain component spectra (solid E_1 , dashed E_2 , dotted E_3): smooth rolloff with no grid-scale shelf; the upwash E_2 amplified above the streamwise E_1 ; grey dotted reference slope $\kappa^{-5/3}$. (b) Migration of the total-spectrum centroid (solid) and the spectral peak (dashed) against the rigid-translation law (grey dotted), with the geometric $\pm\sigma$ ensemble spread shaded: the broadband distortion that linear theory cannot represent, at the operating point and with the noise averaged out.

eddies and viscosity active the component centroids diverge under strain—streamwise migration 1.9 against 1.5 for the upwash at the end of the straining phase, both far below the rigid value 3, the upwash held back most—which is the qualitative departure-from-RDT behaviour CMK report and the spectral signature of the broadband distortion of § 4.3. The comparison is qualitative by construction (initial spectrum, spectral projection and scale separation all differ from the experiment’s); the full account, with both figures and the stated limits, is given in the supplementary material (§ S1), and the quantitative spectral assessment is carried out against the strained-box DNS below, where projection and initial spectrum are matched by construction.

4.6. Validation against a strained-box DNS with an exact-RDT companion

The comparisons above are either model-internal (§ 4.3) or, where external, qualitative by construction (§ 4.5) or at the Reynolds-stress level (§ 3.4). The referent quantity for the leading-edge application is spectral, and the objection that the emergent-spectrum results compare the model only with itself must be met with an independent, projection-matched spectral benchmark under the same strain. To that end we performed direct numerical simulations of plane-strained homogeneous turbulence in a deforming-frame (Rogallo-type) pseudo-spectral formulation: a 128^3 box, decaying isotropic precursor, then plane strain $A = S \text{diag}(\frac{1}{2}, -\frac{1}{2}, 0)$ applied at strain-rate parameters $Sk_t/\varepsilon \in \{0.8, 16\}$ at onset, carried to total strain $e = 1$, four independent realisations ($\kappa_{\max}\eta \approx 0.76$ at $e = 1$: a pilot resolution, so small-scale statements below are qualitative). Alongside every DNS realisation runs an *exact-RDT companion*: the same initial field evolved under the closed-form Cauchy solution of the linearised equations and projected onto lines identically to the DNS and to ODT. The companion is what § 4.2 used to construct the linear reference; here it completes a three-way comparison in which model, linear theory and DNS are reduced to the same one-dimensional observables. On the ODT side the ensembles use an unstrained precursor: the line is evolved as ordinary ODT until the initialisation transient has relaxed and the eddy population is statistically active, and the strain is switched on only then, so that the strained evolution starts from a physically conditioned state rather

$\Delta b_{22}(\kappa_2)$, band κ_2/κ_c :	0.5–0.75	0.75–1.2	1.2–1.9	1.9–3	3–4.8	4.8–7.6	7.6–12
exact projected RDT	+0.146	+0.008	−0.020	−0.041	−0.050	−0.054	−0.056
DNS	+0.032	−0.033	−0.029	−0.032	−0.036	−0.040	−0.046
SC-ODT	+0.103	+0.103	+0.096	+0.073	+0.075	+0.079	+0.076

Table 1. Strain-induced upwash anisotropy per wavenumber band at $e = 1$, $Sk_t/\varepsilon = 0.8$, each system referenced to its own unstrained state. The band below $0.75 \kappa_c$ rests on a few box modes so those results are indicative only.

than from the synthetic initial condition. All ODT statistics in this subsection are medians over 1024 realisations per case with bootstrap confidence intervals; median statistics are used because the strained band energies are strongly intermittent across realisations, and ensemble means of spectral ratios are dominated by rare events at any practical sample size.

4.6.1. Three-way spectral anisotropy: SC-ODT, exact RDT, DNS

The comparison quantity is the strain-induced change of the spectral component anisotropy, $\Delta b_{nn}(\kappa_2) = \phi_{nn}/\sum_m \phi_{mm} - \frac{1}{3}$ at strain e , minus the same quantity evaluated in the system's own unstrained state with each mode mapped to its strained wavenumber. Referencing each system to its own $e = 0$ state cancels viscous decay, which is component-blind, and the longitudinal–transverse structure of the isotropic line spectra, which differs between a three-dimensional box and the component-symmetric ODT line. Figure 11 and Table 1 give the result at $e = 1$, $Sk_t/\varepsilon = 0.8$, on common bands of $\kappa_2(e)/\kappa_c(0)$ with κ_c the initial spectral centroid.

The structure of the tabulated results is the finding. Exact linear theory concentrates the strain-induced upwash anisotropy at the energy-containing scales and removes it from the small scales, because modes aligned with the line are amplified least; the DNS shows the same shape, weakened at the large scales by the nonlinear return—that nonlinear effects absent from linear theory shape the strained anisotropy across scales is likewise the finding of the 4096³ axisymmetric-contraction simulations of Clay & Yeung (2016). SC-ODT distributes its moment-level anisotropy essentially uniformly over wavenumber—the signature, anticipated in § 4.2, of a wavenumber-uniform redistribution operator acting on a uniformly dilated line. The integrated anisotropy is nevertheless correct (the model tracks the closure at the moment level, § 3.4), so the deficiency is one of allocation across scales, not of amount. The transverse splitting $\phi_{11}/\phi_{33} - 1$, a combination of the two transverse components just shown, carries the same reading (not plotted): exact RDT deepens it with wavenumber and the DNS relaxes it towards zero at the small scales, whereas SC-ODT holds it essentially constant—the wavenumber-uniform signature again, now in the neutral-to-streamwise balance rather than the upwash. For the leading-edge application the operative consequence is that the high-wavenumber anisotropy of the upwash spectrum, which the Amiet response weights, has the wrong sign in SC-ODT; the closure element that must carry the correction—a wavenumber-dependent rapid operator in place of the uniform one—is exactly the spectral kernel formulation of § 5, whose introduction is thereby motivated by a measured deficit rather than by its association with an identified physical process.

4.6.2. Distance from linear theory as a function of strain rapidity

The three-way comparison resolves the anisotropy by wavenumber at one strain-rate parameter. A complementary scalar measure tracks the *whole* spectral shape across strain




Figure 11. Three-way comparison at $e = 1$, $Sk_t/\varepsilon = 0.8$: strain-induced anisotropy of the line spectra relative to each system's own unstrained state, resolved by component—(a) Δb_{22} , (b) Δb_{11} , (c) Δb_{33} . SC-ODT (■, 1024 realisations) against exact projected RDT (♦) and DNS (128³, ●, four realisations). The SC-ODT curves are relatively flat in every panel—the wavenumber-uniform signature—whereas exact RDT and the DNS vary with wavenumber.

and strain rapidity: at each total strain the median ODT component line spectra are fitted, jointly in the upwash and transverse components, by the exactly distorted spectrum of linear theory—the closed-form Cauchy distortion of an isotropic von Kármán field, projected onto the line, with only the pre-distortion amplitude and energy wavenumber free. The root-mean-square log-residual of that two-parameter fit is a scale-resolved distance between the model's strained spectra and exact linear theory. Figure 12 shows it for precursor-conditioned ensembles at $Sk_t/\varepsilon \approx 0.4$ and ≈ 8 at onset (1024 realisations each, identical precursor state: the two fits coincide at $e = 0$ at thirteen per cent, the residual shape mismatch of a relaxed ODT line against the von Kármán form).

The distance grows with strain in both cases, but *faster at rapid strain*: to approximately 22 per cent at $e = 2$ for $Sk_t/\varepsilon \approx 0.4$ against approximately 24 per cent, saturating early, at $Sk_t/\varepsilon \approx 8$, with the separation established already by $e = 0.5$. If the model's rapid response were spectrally faithful, the rapid curve is the one that should approach zero: eddy events are subdominant there by construction. That the deviation is largest precisely where the eddies matter least localises the deficiency in the rapid kinematics themselves—the wavenumber-uniform operator and the rigid dilatation—and not in the eddy or relaxation statistics, in quantitative agreement with the three-way comparison. This measured validity envelope accompanies the model into any application: predictions inherit a spectral-shape uncertainty of order the plotted distance at the operating point's strain and rapidity.

5. Closure of the rapid pressure–strain term on a single line

The results of § 4 establish what SC-ODT produces and where it stands against exact linear theory and DNS. The rapid pressure–strain redistribution that drives that evolution was evaluated throughout by the algebraic moment closure of § 2 (Launder *et al.* 1975), applied component by component on the resolved line. That closure is a modelling input, and the redistribution it represents is, in its exact form, a *nonlocal, three-dimensional* object: the rapid pressure at a point is a volume integral over the entire fluctuating field, and the pressure–strain correlation is an integral over the full three-dimensional spectrum—one that can be evaluated scale by scale only where the whole field is available, as Clay & Yeung




Figure 12. Scale-resolved distance between the strained SC-ODT line spectra and exact linear theory: rms log-residual of the two-parameter exact-RDT-distorted von Kármán fit, against total strain, for low and high onset strain-rate parameters (medians over 1024 realisations; identical precursor state at $e = 0$).

(2016) do in a 4096³ simulation of axisymmetric contraction, separating its rapid and slow parts spectrally. What a single one-dimensional line of turbulence can say about such an object is the analytical question of this paper, and this section answers it in a form that is deliberately more modest, and more defensible, than a closure claim: we state exactly what obstructs a single-line representation, we derive sharp bounds on the rapid term from line data alone, we *measure* rather than assume the axisymmetry hypothesis under which the term collapses, and we identify—by measurement—which physical content the moment closure discards.

The section proceeds as follows. The exact term and the obstruction are stated first (§ 5.1–§ 5.2); the axisymmetric collapse follows, together with the statement of what it does *not* deliver (§ 5.3); the bounds that line data do deliver, with their falsifiers, come next (§ 5.4); the axisymmetry assumption is then measured directly in the strained-box DNS, including a control experiment that isolates its failure mode (§ 5.5); the component-ordering deficiency raised by the exact solution is measured and localised (§ 5.6); the reduction to the moment closure and the disjointness from the eddy redistribution close the argument (§ 5.7–§ 5.8).

5.1. The exact rapid pressure–strain term

The starting point is the Poisson equation (2.11) for the fluctuating pressure, whose source splits into a part linear in the fluctuations and the mean velocity gradient (the *rapid* part) and a part quadratic in the fluctuations (the *slow* part) (Pope 2000). The rapid part is linear in $\mathbf{u}'$ and proportional to the mean gradient, which for the present homogeneous plane strain is the constant tensor A_{kl} . Writing p'_r for the rapid pressure and inverting the Laplacian with the free-space Green’s function $G(\mathbf{x}) = -1/4\pi|\mathbf{x}|$,

$$p'_r(\mathbf{x}) = \frac{\rho}{2\pi} A_{kl} \int \frac{1}{|\mathbf{x} - \mathbf{x}'|} \frac{\partial u'_l}{\partial x'_k}(\mathbf{x}') d\mathbf{x}'. \quad (5.1)$$

The rapid pressure–strain correlation that redistributes energy among the Reynolds-stress components is $\Pi_{ij}^{(r)} = \overline{(p'_r/\rho) (\partial u'_i/\partial x_j + \partial u'_j/\partial x_i)}$. Substituting (5.1) and using homogeneity to pass to the velocity spectrum tensor $\Phi_{mn}(\boldsymbol{\kappa})$ (the Fourier transform of $\overline{u'_m(\mathbf{x}) u'_n(\mathbf{x} + \mathbf{r})}$), one obtains the standard exact result (Batchelor & Proudman 1954;

Crow 1968; Hunt & Carruthers 1990)

$$\Pi_{ij}^{(r)} = A_{kl} M_{ijkl}, \quad M_{ijkl} = \int_{\mathbb{R}^3} \left[\frac{\kappa_j \kappa_k}{\kappa^2} \Phi_{il}(\boldsymbol{\kappa}) + \frac{\kappa_i \kappa_k}{\kappa^2} \Phi_{jl}(\boldsymbol{\kappa}) \right] d\boldsymbol{\kappa}, \quad (5.2)$$

where $\kappa = |\boldsymbol{\kappa}|$. Equations (2.11) and (5.2) are exact; no modelling has entered. The tensor M_{ijkl} is a linear functional of the *full three-dimensional* spectrum $\Phi_{mn}(\boldsymbol{\kappa})$ through the nonlocal projection operator $\kappa_a \kappa_b / \kappa^2$, which is the wavenumber-space image of the inverse Laplacian in (5.1). This nonlocality and three-dimensionality are the obstruction that the single-line model must confront.

5.2. The obstruction: what an ODT line resolves

In the present formulation the line carries the fluctuation velocity $\mathbf{u}'(\mathbf{y})$ about a prescribed uniform mean gradient—standard ODT can also evolve a mean profile—on a single line aligned with the resolved direction, here taken along the contraction axis of the strain, $y \equiv x_2$. From it the model has access to the one-dimensional spectra

$$\phi_{mn}(\kappa_2) = \int_{\mathbb{R}^2} \Phi_{mn}(\boldsymbol{\kappa}) d\kappa_1 d\kappa_3, \quad (5.3)$$

obtained by integrating Φ_{mn} over the two unresolved wavenumber components. The integrand of (5.2), by contrast, carries the *direction-resolved* weight $\kappa_a \kappa_b / \kappa^2$, which cannot be reconstructed from $\phi_{mn}(\kappa_2)$ alone: integrating out κ_1, κ_3 as in (5.3) destroys exactly the angular information that the projection operator acts on. The reduction $M_{ijkl} \rightarrow$ (functional of ϕ_{mn}) is therefore not available without an assumption about how the energy is distributed over the unresolved wavevector directions. This is the crux of the present question, and we make the required assumption explicit rather than implicit.

The obstruction just stated does not disappear under any symmetry assumption short of full isotropy. As § 5.3 shows, axisymmetry about the line reduces the unknown spectrum to two scalar functions of $(\kappa_2, \kappa_\perp)$ —but the line data are two weighted $\kappa_\perp$ -integrals of those functions per resolved wavenumber, so the map from what the line carries to what the rapid term requires retains an infinite-dimensional null space. Stating this null space exactly, and extracting what survives it, is the correct formulation of the single-line question.

5.3. The closure assumption and the collapsed form

The spectrum tensor of the incoming turbulence, before the rapid distortion acts, is axisymmetric about the resolved direction $\mathbf{e}_2$: $\Phi_{mn}(\boldsymbol{\kappa})$ depends on $\boldsymbol{\kappa}$ only through κ_2 and $\kappa_\perp = (\kappa_1^2 + \kappa_3^2)^{1/2}$, and its polarisation is invariant under rotation about $\mathbf{e}_2$.

This is a statement about the *turbulence*, not about the strain. It is justified for the present configuration because the inflow is grid turbulence, which is closely isotropic at the grid and is subsequently contracted along a single axis as it approaches the stagnation plane; isotropy is the special case of assumption A1 (see Appendix A.1), and uniaxial contraction preserves axisymmetry about the contraction axis. The plane strain $A_{kl} = \text{diag}(\frac{1}{2}, -\frac{1}{2}, 0)$ acting at the leading edge is itself *not* axisymmetric, so A1 is imposed on the state on which the rapid operator acts, consistent with the rapid-distortion ordering in which $\Pi^{(r)}$ is evaluated on the current (here, axisymmetric to leading order) field. The error incurred when the accumulated distortion has driven the spectrum away from axisymmetry is quantified in § 5.5.

Under A1 the angular integral in (5.2) can be carried out. Writing $\boldsymbol{\kappa} = (\kappa_\perp \cos \theta, \kappa_2, \kappa_\perp \sin \theta)$ and using that, by A1, Φ_{mn} is independent of θ , the azimuthal averages of the projection

weights reduce to

$$\left\langle \frac{\kappa_1^2}{\kappa^2} \right\rangle_\theta = \left\langle \frac{\kappa_3^2}{\kappa^2} \right\rangle_\theta = \frac{1}{2} \frac{\kappa_\perp^2}{\kappa^2}, \quad \left\langle \frac{\kappa_2^2}{\kappa^2} \right\rangle_\theta = \frac{\kappa_2^2}{\kappa^2}, \quad \left\langle \frac{\kappa_1 \kappa_3}{\kappa^2} \right\rangle_\theta = 0, \quad (5.4)$$

with $\kappa^2 = \kappa_2^2 + \kappa_\perp^2$, and the weights mixing a resolved and an unresolved index (e.g. $\kappa_1 \kappa_2 / \kappa^2$) averaging to zero. Substitution of (5.4) into (5.2) removes the azimuthal dependence and leaves a two-dimensional integral over $(\kappa_2, \kappa_\perp)$ of the axisymmetric spectrum. Defining the diagonal one-dimensional spectra along the line,

$$E_n^\parallel(\kappa_2) = \int_0^\infty \Phi_{nn}(\kappa_2, \kappa_\perp) 2\pi \kappa_\perp d\kappa_\perp \quad (\text{no sum}), \quad (5.5)$$

the rapid pressure–strain term reduces to a two-dimensional integral over the half-plane $(\kappa_2, \kappa_\perp)$ of the axisymmetric spectrum. The axisymmetric, solenoidal spectrum is fixed by two scalar functions of $(\kappa_2, \kappa_\perp)$ (Batchelor 1946; Chandrasekhar 1950); we denote them $a(\kappa_2, \kappa_\perp)$ and $c(\kappa_2, \kappa_\perp)$, where a is the isotropic-like scalar (the sole survivor when the turbulence is isotropic) and c measures the axisymmetric departure from isotropy. Carrying out the azimuthal average (5.4) and specialising to the plane strain $A = \text{diag}(\frac{1}{2}, -\frac{1}{2}, 0)$, the diagonal rapid pressure–strain components are

$$\Pi_{nn}^{(r)} = \int_0^\infty \int_0^\infty g_n(x) [a(\kappa_2, \kappa_\perp), c(\kappa_2, \kappa_\perp)] 2\pi \kappa_\perp d\kappa_\perp d\kappa_2 \quad (\text{no sum on } n), \quad (5.6)$$

with $x \equiv \kappa_2^2 / \kappa^2$, $\kappa^2 = \kappa_2^2 + \kappa_\perp^2$, and the kernels, linear in the two spectral scalars,

$$g_1(x) = a\left(\frac{1}{8} + \frac{3}{4}x - \frac{7}{8}x^2\right) + c\frac{7}{8}x(1-x)^2, \quad (5.7)$$

$$g_2(x) = -\frac{3}{2}x [a(1-x) + c(1-x)^2], \quad (5.8)$$

$$g_3(x) = a\left(-\frac{1}{8} + \frac{3}{4}x - \frac{5}{8}x^2\right) + c\frac{5}{8}x(1-x)^2. \quad (5.9)$$

Equations (5.6)–(5.9) constitute the axisymmetric collapse: *under A1, the rapid pressure–strain term is a functional of the axisymmetric spectral scalars a, c on the resolved half-plane and the mean strain*, with no reference to the unresolved azimuthal structure beyond what A1 fixes. The kernels satisfy two exact constraints that verify the reduction. First, energy conservation: pressure–strain redistributes without changing the trace, and indeed $g_1(x) + g_2(x) + g_3(x) = 0$ identically for all x and for both scalars, so $\Pi_{11}^{(r)} + \Pi_{22}^{(r)} + \Pi_{33}^{(r)} = 0$. Second, the isotropic limit: setting $c = 0$ with a a function of κ alone, the integral of g_3 over the half-plane vanishes, so $\Pi_{33}^{(r)} = 0$ —as it must, since the plane strain has no component along x_3 and the isotropic rapid term is proportional to S_{ij} (Crow 1968). Both constraints hold identically in the derived kernels, confirming the tensor reduction.

The status of (5.6)–(5.9) must be stated precisely, because it is here that the closure question is decided. Under A1 the rapid term is a linear functional of the two axisymmetric scalars $a(\kappa_2, \kappa_\perp)$ and $c(\kappa_2, \kappa_\perp)$ on the resolved half-plane—a *two-dimensional* representation, not yet a single-line closure, since the line determines only the $\kappa_\perp$ -integrals (5.5) of fixed combinations of a and c and not the functions themselves. Equations (5.6)–(5.9) are therefore the collapsed form on which any single-line statement must be built. What the line data actually determine—and how sharply—is established next.

constraint level	$\mathcal{I}_1$	$\mathcal{I}_2$	$\mathcal{I}_3$
line data + solenoidality + realisability	0.77	1.28	0.86
+ log-Lipschitz slope cap ($\lambda = 4$)	< 5% narrower		
+ polarisation cap $ c \leq \gamma a$, $\gamma = 0$ (i.e. $c = 0$)	0.34	0.58	0.24

Table 2. Kinematic indeterminacy $\mathcal{I}_n = (\Pi_{nn}^+ - \Pi_{nn}^-)/2k_t S$ of the three rapid pressure–strain components under successively tighter side conditions, calibrated on an isotropic von Kármán spectrum. Line data with solenoidality and realisability leave an $O(1)$ indeterminacy; the log-Lipschitz smoothness cap barely narrows it, whereas the polarisation cap ($\gamma = 0$, i.e. $c = 0$) roughly halves it.

5.4. What line data determine: sharp bounds on the rapid term

Under A1 the solenoidal spectrum tensor takes the two-scalar form $\Phi_{mn} = a P_{mn} + c \xi_m \xi_n$, with P_{mn} the solenoidal projector and ξ_m the projected axis direction; realisability is the pair of pointwise conditions $a \geq 0$ and $a + (1 - x)c \geq 0$ with $x = \kappa_2^2/\kappa^2$. The line carries $\phi_{11}(\kappa_2) = \phi_{33}(\kappa_2)$ and $\phi_{22}(\kappa_2)$: per resolved wavenumber, two weighted $\kappa_\perp$ -integrals of the two unknown functions. The rapid term (5.6) is linear in (a, c) , so its extreme values over everything consistent with the line data, solenoidality and realisability are obtained from a family of linear programs, one per κ_2 , yielding sharp bounds $[\Pi_{nn}^-, \Pi_{nn}^+]$ and a kinematic indeterminacy $\mathcal{I}_n = (\Pi_{nn}^+ - \Pi_{nn}^-)/2k_t S$ that quantifies exactly how much the null space costs. The forward map from (a, c) to the line spectra and the derivation of the bound kernels are given in Appendix A.

Calibrated on an isotropic von Kármán spectrum (with a verification suite confirming tracelessness, the isotropic angular averages $+\frac{1}{5} : -\frac{1}{5} : 0$, the longitudinal–transverse relation, and the exact $\frac{4}{5}k_t S_{ij}$), the indeterminacy is

Two features of the Table 2 matter. Smoothness assumptions buy almost nothing: the null space is not a roughness artefact. The polarisation cap is the lever—but § 5.5 shows that under strain the effective polarisation reaches order one, so the honest bound in the strained regime is the uncapped one.

As shown in Appendix A, the representation also supplies two *free falsifiers* that the line itself reports: under A1, $\phi_{11}(\kappa_2) = \phi_{33}(\kappa_2)$ at every wavenumber (T1) and the velocity cross-spectra vanish (T2). Any measured transverse splitting is a wavenumber-resolved violation of A1; rather than assuming the hypothesis, the line monitors it.

5.5. The axisymmetry hypothesis, measured

A1 is exact for the isotropic inflow and is degraded by the plane strain itself. Rather than estimating this error perturbatively, we measure it in the strained-box DNS of § 4.6, where the full spectrum is available: the exact azimuthal decomposition of Φ_{22} yields the true scalars (a, c) , the discarded $m = 2$ azimuthal residue (the direct A1-violation measure, with isotropic sampling floor 0.133 in this box), and the effective polarisation $\gamma_{\text{eff}} = |c|/a$.

Three conclusions follow (table 3). First, the A1 error is *not* perturbatively small in the accumulated anisotropy: the residue grows from its floor to 0.19 at $Sk_t/\varepsilon = 0.8$ and to 0.40 at 16 by $e = 1$ —larger at rapid strain, because at moderate strain rate the nonlinear dynamics restore axisymmetry as fast as the strain destroys it. An error estimate quadratic in the anisotropy ($O(b^2)$), which the symmetric structure of the discarded harmonic suggests, is superseded by this measurement. Second, the polarisation cap that narrows the isotropic bounds is illegitimate under strain: γ_{eff} passes unity by $e \simeq 0.5$. Third, correspondingly,

e	$Sk_t/\varepsilon = 0.8$					$Sk_t/\varepsilon = 16$				
	b_{22}	γ_{eff}	$m=2$	$\frac{\phi_{11}}{\phi_{33}}-1$		b_{22}	γ_{eff}	$m=2$	$\frac{\phi_{11}}{\phi_{33}}-1$	
0	0.003	0.10	0.133	+0.06		0.003	0.10	0.133	+0.06	
0.5	0.064	0.79	0.166	-0.20		0.070	1.28	0.247	-0.21	
1	0.125	1.44	0.189	-0.32		0.124	4.22	0.396	-0.46	

Table 3. A1 measured under plane strain (128³ DNS, four realisations): upwash anisotropy, effective polarisation, azimuthal $m = 2$ residue of Φ_{22} (isotropic floor 0.133), and the on-line falsifier T1.




Figure 13. The A1 control experiment at $Sk_t/\varepsilon = 16$ (128³, four realisations): plane strain (A1 broken; solid, filled circles) against axisymmetric contraction at matched A_{22} (A1 exact by symmetry; dashed, open squares). (a) Azimuthal $m = 2$ residue of Φ_{22} : it grows under plane strain but stays at the isotropic floor (grey dotted) under axisymmetric strain. (b) Transverse splitting $\phi_{11}/\phi_{33} - 1$ for each geometry, and the upwash anisotropy b_{22} (grey dash-dot, indistinguishable between geometries)—the azimuthal structure that A1 assumes away is what differs, while the acoustically relevant upwash amplification is robust.

the uncapped bound brackets the exact rapid term (evaluated by direct integration in the DNS) up to $e \approx 0.25$ – 0.5 and fails beyond, and the failure is A1's, not the estimator's.

The decisive evidence for that attribution is a control experiment: the same box under *axisymmetric contraction* with matched A_{22} , for which A1 is exact by symmetry. There the $m = 2$ residue stays at the sampling floor and the transverse splitting at zero at every strain up to $e = 1$ even at $Sk_t/\varepsilon = 16$, while under plane strain both grow (figure 13). The single-line closure is therefore *exact for axisymmetric distortion*; its plane-strain error is a strain-geometry property, now quantified, rather than a defect of the line representation. For the applications this maps cleanly: a rounded, three-dimensional stagnation region imposes an axisymmetric strain (closure exact); a two-dimensional leading edge imposes plane strain (closure carries the measured correction). The acoustically relevant amplification $b_{22}(e)$, finally, is indistinguishable between the two geometries (0.124 against 0.117 at $e = 1$, $Sk_t/\varepsilon = 16$): the quantity the leading-edge response consumes is robust to precisely the geometry that A1 is sensitive to.

5.6. Component ordering under plane strain

The exact solution of the rapid problem under plane strain has the neutral spanwise component eventually overtaking the upwash; algebraic closures predict the opposite. The

exact rapid term evaluated in the DNS settles the matter: at $e = 1$, $\Pi_{33}^{(r)}/k_t S = 0.14$ at $Sk_t/\varepsilon = 0.8$ and 0.27 at 16 , exceeding $\Pi_{11}^{(r)} = 0.23$ in the rapid case—the overtaking is real and sets in at rapid strain. LRR, whose spanwise component is carried by its b -dependent term, matches at $Sk_t/\varepsilon = 0.8$ ($\Pi_{33}^{(r)} = 0.143 k_t S$) but delivers half the exact value at 16 ; IP gives zero identically. The deficiency thus lives in the moment closure’s b -term, not in the line model: the ODT moment-level b_{ij} follows the closure to $\sim 2\%$, and its b_{33} stays within 0.006 of zero in every run. In the spectral representation the ordering is carried by the c -scalar through the kernels (5.7)–(5.9); the moment reduction omits it. This answers, by measurement, where the known component-ordering error enters and what it would take to remove it.

5.7. Reduction to the moment closure

The closed form (5.6)–(5.9) operates on the resolved spectrum, whereas the distortion results of § 4 were obtained with the algebraic moment closure of § 2.4 (Launder *et al.* 1975). For the two to be consistent, the spectral kernel must reduce to that moment closure in the limit where the moment closure is exact—isotropic turbulence. We show that it does, which both justifies the closure used in the simulations and fixes the prefactor in an explicit convention (the coefficient of (5.12) and the remark following it).

Figure 14 plots the kernels g_1, g_2, g_3 of (5.7)–(5.9) against the angular variable $x = \kappa^2/\kappa^2$. Two features are exact and hold for both spectral scalars. The kernels sum to zero at every x ,

$$g_1(x) + g_2(x) + g_3(x) = 0 \quad \forall x, \quad (5.10)$$

so the rapid term conserves turbulent kinetic energy pointwise in wavenumber, not merely in the integral; and each kernel vanishes at $x = 1$, where the wavevector aligns with the line and the projection $\kappa_a \kappa_b / \kappa^2$ degenerates. The signs are physically ordered: $g_1 > 0$ (energy fed to the streamwise component E_1), $g_2 < 0$ (energy removed from the line/upwash component E_2) over the bulk of the angular range, with g_3 changing sign, consistent with the plane strain stretching x_1 and compressing x_2 .

For *isotropic* turbulence the anisotropy scalar vanishes ($c = 0$) and the isotropic-like scalar a depends only on κ . Averaging the kernels over solid angle—equivalently, integrating $g_n(x)$ against the isotropic distribution of x on the sphere—gives

$$\langle g_1 \rangle : \langle g_2 \rangle : \langle g_3 \rangle = +\frac{1}{5} : -\frac{1}{5} : 0, \quad (5.11)$$

in the ratio $+1 : -1 : 0$. This is exactly the structure of the plane strain $S_{ij} = \text{diag}(\frac{1}{2}, -\frac{1}{2}, 0)$: the rapid term is proportional to S_{ij} , with $\Pi_{33}^{(r)} = 0$ because the strain has no component along x_3 . The spectral kernel therefore collapses, under isotropy, to the isotropic-production form of the moment closure,

$$\Pi_{ij}^{(r)} = \frac{4}{5} k_t S_{ij} \quad (\text{isotropic limit}), \quad (5.12)$$

the result of Crow (1968) that underlies the isotropisation-of-production (IP) rapid term in the closure of Launder *et al.* (1975). We note a convention: with the symmetrised definition of the pressure–strain correlation adopted here and the factor of two carried in the Poisson source (2.11), the kernels (5.7)–(5.9) yield the coefficient $\frac{2}{5} k_t$ for the 11-component directly (since $S_{11} = \frac{1}{2}$), equivalent to the standard $\frac{4}{5} k_t S_{ij}$ of (5.12); the ratio (5.11) and the tracelessness (5.10) are independent of this convention.

The consequence is that the algebraic closure used to obtain the distortion of § 4 is not an independent assumption but the isotropic-limit reduction of the line-consistent spectral




Figure 14. Rapid pressure–strain closure kernels g_1 (circles), g_2 (squares), g_3 (triangles) (eqs. (5.7)–(5.9)) against the angular variable $x = \kappa_2^2 / \kappa^2$. Solid curves: coefficient of the isotropic-like spectral scalar a ; faint grey dashed curves: coefficient of the anisotropy scalar c . The kernels sum to zero at every x (energy conservation, eq. (5.10)) and vanish at $x = 1$. Averaged over an isotropic spectrum they give the ratio $+\frac{1}{5} : -\frac{1}{5} : 0$ (eq. (5.11)), recovering the isotropisation-of-production form $\Pi_{ij}^{(r)} = \frac{4}{5} k_t S_{ij}$ (Crow 1968) that underlies the moment closure used in § 4.

term derived here. The spectral form (5.6) carries, in addition, the angular (x -dependent) and anisotropy (c -scalar) information that the moment closure omits—and the effects of both omissions are now measured: the wavenumber dependence of the linear response in the three-way comparison of § 4.6.1, and the component ordering in § 5.6. The moment closure and the collapsed form thus agree to leading order, and the terms in which they differ are exactly the measured deficiencies of § 4.

5.8. Consistency with the eddy redistribution: avoiding double counting

A closed rapid term is necessary but not sufficient: it must not duplicate redistribution already effected by the eddy events. The pressure–strain correlation decomposes canonically into rapid and slow parts, $\Pi_{ij} = \Pi_{ij}^{(r)} + \Pi_{ij}^{(s)}$, corresponding to the two source terms in (2.11). The slow part $\Pi_{ij}^{(s)}$ is the return-to-isotropy mechanism, quadratic in the fluctuations and independent of the mean gradient; the eddy events in ODT, which act on the fluctuating field without reference to A_{kl} , are the model’s representation of exactly this slow redistribution. The rapid part $\Pi_{ij}^{(r)}$, proportional to A_{kl} and linear in the fluctuations, is by construction absent from the eddy mechanism, which carries no dependence on the mean strain. The two therefore occupy disjoint terms of the canonical decomposition, and adding (5.6) as an explicit mean-strain-driven source alongside the eddies does not double-count, provided the eddy events are calibrated to reproduce the slow return-to-isotropy in the absence of mean strain—a condition that the no-strain baseline of § 4.3 verifies directly, since there the centroid relaxes under the eddies alone with no spurious rapid contribution. This argument establishes consistency at the level of the decomposition; the no-strain baseline of § 4.3 provides the direct numerical confirmation, the centroid there relaxing under the eddy mechanism alone.

The disjointness argument is now also a measured statement, from both sides. With the strain switched off the eddy events generate no spurious strain-like signal (the nominal-turbulence null of Appendix B). Conversely, in the rapid limit ($Sk_t/\varepsilon = 16$, under one eddy per line in $e \leq 1$) every kernel-energy allocation rule tested—the isotropic kernel, a strain-biased generalisation, and discrete eddy types—reproduces the exact rapid-distortion evolution of $b_{22}(e)$ to ± 0.003 : the rapid term cannot be reached from inside the eddy events, and the continuous operator carries it alone. Rapid and slow redistribution occupy disjoint terms by construction, and now by measurement.

5.9. Summary: what the single-line treatment establishes

1. The rapid term is exactly $A_{kl}M_{ijkl}$; its single-line representation is obstructed by an azimuthal null space that is stated exactly and survives the axisymmetric collapse at the level of the two spectral scalars.
2. Under A1 the term collapses to the kernel form (5.6)–(5.9); A1 is exact for axisymmetric distortion and is *measured* under plane strain, where its violation grows with accumulated strain and with strain rapidity (table 3).
3. Independently of A1’s validity DNS, line data bound the rapid term sharply; the bound is valid—and verified to bracket the exact term—up to $e \approx 0.25$ – 0.5 under plane strain, and at all computed strains under axisymmetric distortion.
4. The moment closure used in the simulations is the isotropic reduction of the collapsed form (verified, including the exact $\frac{4}{3}k_t S_{ij}$); what it discards is the wavenumber dependence of the linear response and the component ordering, both measured in § 4 and § 5.6.
5. Rapid and slow redistribution are disjoint: verified by the strain-off null and by the rapid-limit invariance of every tested kernel-energy allocation rule.
6. The acoustically relevant upwash amplification $b_{22}(e)$ is linear to $e = 1$ and robust to the strain geometry that controls A1’s validity.

6. Measured limits, and scale-local relaxation

The envelope of § 4.6.2 localises the model’s spectral deficiency in its wavenumber-uniform rapid kinematics: the strain-induced anisotropy is allocated uniformly across scales while linear theory and DNS concentrate it at the energy-containing range and remove it from the fine scales. This section asks the constructive question: which mechanism available to the model can supply the missing scale dependence? We answer it experimentally, by a scale-resolved diagnostic and a controlled family of interventions in the eddy mechanism, each tested at 1024 realisations with seeds paired to a common baseline, and each either falsified or confirmed by the same measurement. The outcome is a single measured statement: what is missing is *scale-local relaxation*, and supplying it—at the scales of the imbalance—is both necessary and, at moderate strain, sufficient to move the fine-scale anisotropy.

6.1. A transmission diagnostic

The whole investigation is carried out at a single, moderate turbulence Reynolds number, chosen so that the diagnostic has room to separate scales but the cascade is short: the molecular viscosity ($\nu = 10^{-4}$; Appendix B) is set so that more than a decade of scales is resolved below the energy peak with wide margin. This regime is the referent of the “at this

Reynolds number” qualifications that follow, and the caveat is real—the scale separation, not the eddy statistics, is what the negative results below turn on. The configuration is a scale-diagnostic variant of the strained line: unit strain rate (so $e = t$), an initial spectrum peaked at a wavenumber κ_p with the resolved decade of smaller scales below it, and two measurement bands in wavenumber κ_2 : a low-wavenumber, energy-containing band $\kappa_2 \in [0.6, 2] \kappa_p$ (subscript lo) and a high-wavenumber, fine-scale band $[6, 16] \kappa_p$ (subscript hi)—the latter roughly two triplet-map generations *finer*, since each map compresses scales by a factor of three. On each band the component imbalance is $\mathcal{A}_b = E_2^\parallel(b)/\frac{1}{2}[E_1^\parallel(b) + E_3^\parallel(b)]$, the band-integrated upwash energy over the transverse mean; all statistics are medians over realisations with bootstrap confidence intervals, for the intermittency reason given in § 4.6. The model’s eddy and kernel mechanisms are component-symmetric, so $\mathcal{A} = 1$ is their fixed point; the strain drives $\mathcal{A}$ up at the scales where it deposits anisotropy, and the ratio $\mathcal{A}_{\text{hi}}/\mathcal{A}_{\text{lo}}$ measures how much of the large-scale imbalance survives transmission down the cascade.

SC-ODT (C0.2; supplementary material, figure S3, black; also the black curve of figure 15) exhibits the deficiency of § 4.6.1 in transmission form: $\mathcal{A}_{\text{hi}}/\mathcal{A}_{\text{lo}}$ stays near unity out to $e = 3$ (1.02, 0.95, 0.89 at $e = 1, 2, 3$)—the anisotropy injected at the energy-containing scales crosses two map generations essentially undiminished—and drops only at the deepest strain. The mechanism is direct: the high band is populated by the maps of *large* eddies (one generation takes κ_p to $3\kappa_p$, a second to $9\kappa_p$, inside the band), so fine-scale anisotropy is delivered by large-eddy events, not by a chain of small ones.

6.2. Interventions: what fails and what works

The mechanism was located by a controlled search: every intervention category—acceptance gating (C1), map-image geometry (C2), sub-scale relaxation kernels (C3), concurrent relaxation (C4)—was run at 1024 paired-seed realisations against SC-ODT (C0.2), and each variant’s knob, measured fine-scale effect and remarks are tabulated in the supplementary material (§ S2, table S1 and figure S3; the category.variant labels C0–C5 are defined with the reference configurations in Appendix B). The search returns one positive result and one shared limitation, documented case by case in the supplementary material and summarised here. Gating and map-image geometry are null where it matters: the gates modulate the eddy rate, not the scale allocation, and removing up to 97 per cent of small-eddy events leaves the fine-scale anisotropy and energy unchanged—the measurement band is fed directly by large-eddy maps, so no acceptance-gating scheme can control transmission at this Reynolds number. What works is relaxation applied *at the scale of the imbalance*: recursing the energy-equalising kernel procedure of § 2 into the map images, at the depths that reach the measured band ($\ell/9$ and $\ell/27$; C3.2, C3.3), produces the first statistically significant fine-scale isotropisation of the search ($\mathcal{A}_{\text{hi}}$ from 1.43 to 1.27 at $e = 1$) with band energies, the upwash energy fraction and the intermittency tails untouched. Every eddy-locked refinement of that success—depth tied to eddy size (C3.4), per-event iteration (C3.5), delocalisation and stacking (C4)—then shares a single limitation: the fine-scale gains fade beyond $e \approx 3$. Depth sets *where* the relaxation acts and is necessary; dosage and delocalisation raise the early gains in proportion to the total relaxation delivered; and no eddy-locked combination—more depth, more dosage, more events, other locations—prevents the fade. The remaining suspect is the clock itself: relaxation acts at the eddy rate while the strain resupplies the imbalance continuously. This motivates the final intervention.

6.3. Decoupling the relaxation clock

The one feature the eddy-locked formulations do not modify is *when* relaxation acts. We therefore give it its own clock. Whereas the eddy events are sampled from the state-dependent rate density $\lambda(y_0, \ell; t)$ of § 2.1, the relaxation-only events form an independent Poisson process with the prescribed rate density

$$\lambda_c(y_0, \ell) = \frac{R f(\ell)}{L - \ell}, \quad f(\ell) = \frac{2L_p \ell^{-2} e^{-2L_p/\ell}}{e^{-2L_p/L_{\max}} - e^{-2L_p/L_{\min}}}, \quad \ell \in [L_{\min}, L_{\max}], \quad (6.1)$$

with R the total event rate, the placement y_0 uniform over the $L - \ell$ admissible positions on the domain of length L , and f the presumed candidate-size density of the model's own eddy sampling (Kerstein 1999; Stephens & Lignell 2021), its most-probable size L_p and range $[L_{\min}, L_{\max}]$ taken from the SC-ODT run—so no empirically gathered size statistics enter and, unlike λ , (6.1) carries no dependence on the flow state. This candidate density is the model's presumed *sampling* distribution, not the realised distribution of accepted eddies, which the acceptance rule reshapes; it is used because it is convenient to sample and carries the gross features of the physical eddy-size distribution. Each event applies the kernel of (2.2) to its interval—the interval's own triplet-map displacement as the kernel shape, no map applied, amplitudes (2.5) from the interval's own integrals—followed by the adaptive-depth hierarchy of § 6.2 inside it; the first identity of (2.3) conserves each momentum component and the transfer matrix of (2.4) their summed kinetic energy, exactly and per event. The mean rate of the relaxation events is matched to the SC-ODT mean eddy rate, $R = \bar{N}_e/T$ with $\bar{N}_e$ the mean number of eddy events over an SC-ODT realisation of duration T , so the comparison against the eddy-locked variants is dose-matched by construction.

The result (figure 15) decides the question. With the clock, the fine-scale imbalance is reduced significantly at every strain and the effect grows rather than fades: the paired contrasts reach -0.30 in the low band and remain significant in the high band at $e = 3.9$. A constant-rate event sequence keeps pace precisely where the eddy rate could not, confirming that event-locked timing—not any per-event property—was the deep-strain limiter.

However, an event sequence that draws from the full candidate density (capped only at the model's default largest scale $L_{\max}$) is scale-blind and over-corrects. It also suppresses the one-point anisotropy that the strain should produce, pulling $u_2^2/2k_t$ from 0.58 to 0.41 at $e = 3.9$ and thereby breaking the moment-level behaviour validated in § 3.4.

Restricting the event sequence to sub-band interval sizes—replacing f in (6.1) by its restriction to $\ell \leq \ell^*$, renormalised (over-cap draws are resampled), where ℓ^* separates the two measurement bands—resolves the tension. The size-capped clock (the “clock-relaxed ODT” of the figures) holds the fine scales at every strain with no fade ($\mathcal{A}_{\text{hi}} = 1.30\text{--}1.44$ against SC-ODT's 1.63–1.94), with contrasts that remain significant throughout and are strongest at $e = 3$ (figure 15a). The one-point trajectory is essentially untouched in the moderate-strain regime ($u_2^2/2k_t$ within 0.005 of SC-ODT at $e = 1$) and only mildly reduced at the deepest strain (figure 15b).

The mean rate of relaxation events is a single control: quadrupling the rate extends the upward trend in figure 15a at a proportionate cost in figure 15b.

6.4. The measured missing ingredient

The deficiency measured by the three-way comparison—anisotropy allocated uniformly over wavenumber—is confirmed by the transmission diagnostic of § 6.1 in a different observable: the band-imbalance ratio $\mathcal{A}_{\text{hi}}/\mathcal{A}_{\text{lo}}$ stays near unity, so the strain-induced im-




Figure 15. The independent relaxation clock (C5) against the eddy-locked categories. (a) Fine-scale-band imbalance $\mathcal{A}_{hi}$ against total strain: the eddy-locked variants fade beyond $e \approx 3$; the clock holds at every strain. (b) The price: upwash energy fraction. The comparison is the SC-ODT (C0.2) trajectory—the moment-level behaviour validated in § 3.4—which the uncapped clock suppresses; the size cap at matched rate preserves it, and the $4 \times$ rate (C5.3) reduces it in proportion. Medians over 1024 paired-seed realisations; the clock rate is matched to the SC-ODT eddy rate.

balance reaches the fine scales undiminished. It is unfixable by gating or map reweighting, and is cured by exactly two ingredients, each isolated by a falsified alternative: relaxation applied *at the scale of the imbalance* (depth, not dosage or location) and relaxation applied *on its own clock*, with the event-size distribution capped at $\ell \leq \ell^*$ (the clock-relaxed ODT of § 6.3), which alone survives deep strain. One limit remains, and it is structural rather than curable by any amount of relaxation: an isotropising mechanism is bounded at component equality, whereas exact linear theory and the DNS carry the fine-scale upwash *below* the transverse mean. Crossing that boundary requires the rapid operator itself to discriminate by wavenumber, and the formulation already contains its natural form: the spectral kernel (5.6)–(5.9) is precisely a wavenumber- and angle-dependent rapid operator, whose discretisation in place of the wavenumber-uniform moment closure of § 2 is the identified next step.

The comparison that summarises this part is figure 16: the strain-induced upwash anisotropy per wavenumber band—the observable on which § 4.6.1 established the benchmark—for exact linear theory, the DNS, SC-ODT, and the model under each of the two mechanism forms (adaptive depth; the size-capped clock), at three strain-rate parameters, from same-binary ensembles under the identical precursor protocol. Four statements can be read directly from the shared axes. First, at $Sk_t/\varepsilon = 16$ the DNS lies on the exact-RDT curve: the rapid limit is realised in the benchmark itself, and the deficiency of SC-ODT—its uniform allocation of a correct integrated anisotropy—is starkest precisely where the eddies matter least. Second, at moderate rapidity the mechanism tilts the model’s allocation towards the DNS—the anisotropy now decreases with wavenumber, the fine-band excess roughly halved (at $Sk_t/\varepsilon = 0.8$: from +0.076 against the DNS value -0.046 to +0.024)—and the improvement grows as the rapidity falls, the trend the application regime sits on. Third, within $e \leq 1$ the two mechanism forms are equivalent on these axes; the clock’s advantage is the deep-strain hold of figure 15, which this protocol—bounded at $e = 1$ by the DNS reference—cannot display. Fourth, at $Sk_t/\varepsilon = 16$ neither form moves the allocation: rapid




Figure 16. The five-way comparison on the benchmark observable of § 4.6.1: strain-induced upwash anisotropy per wavenumber band at $e = 1$, for exact linear RDT, DNS, SC-ODT, and SC-ODT with each mechanism form—scale-local relaxation at adaptive depth, and the clock-relaxed variant (size-capped relaxation clock)—at $Sk_t/\varepsilon = 0.4, 0.8$ and 16 (ODT: 1024 realisations per case, same binary and precursor protocol; DNS and companions as in § 4.6; no DNS exists at 0.4 , where exact RDT is the reference—at fixed total strain it is independent of the strain rate). At rapid strain the DNS realises the exact-RDT allocation; the mechanism moves the model towards the DNS increasingly as the rapidity falls.

1013 strain limits the physical time available, and no relaxation scheme manufactures more of
 1014 it. The remaining, sign-level gap belongs to the wavenumber- and angle-dependent rapid
 1015 operator of (5.6)–(5.9), identified above as the next step.
 1016

1017 7. Consequence for leading-edge noise

The application that motivates the model can now be served with its output. In the leading-edge-noise theory of Amiet (1975) the far-field pressure spectral density of a large-aspect-ratio aerofoil factorises as a product of the flow and observer geometry, the aerofoil response $|\mathcal{L}|^2(K_x)$, the upwash wavenumber spectrum $\Phi_{ww}(K_x, 0)$ and the spanwise correlation length $\ell_y(K_x) = \pi \Phi_{ww}(K_x, 0) / \int \Phi_{ww}(K_x, K_y) dK_y$, with K_x the streamwise and K_y the spanwise wavenumber. For a *change of inflow spectrum* at fixed aerofoil, flow and observer, everything except the turbulence kernel

$$\mathcal{K}(K_x) = \pi \frac{\Phi_{ww}(K_x, 0)^2}{\int \Phi_{ww}(K_x, K_y) dK_y} \quad (7.1)$$

1018 cancels exactly, so the change of radiated level is $\Delta \text{SPL}(K_x) = 10 \log_{10}[\mathcal{K}^{\text{distorted}}/\mathcal{K}^{\text{frozen}}]$,
 1019 independent of the aerofoil geometry and response. This is the sharpest question the
 1020 distortion model can be asked without importing an acoustic solver: by how much does the
 1021 prediction change when the inflow spectrum is the *computed, distorted* one instead of the
 1022 frozen upstream spectrum that standard practice prescribes?

1023 The distorted input is taken from the model's own strained ensembles in the representa-
 1024 tion already fitted in § 4.6.2: the exact-RDT-distorted von Kármán family, two parameters
 1025 (A_0, κ_e) per strain, fitted to the median line spectra of the SC-ODT 1024-realisation
 1026 ensembles at $Sk_t/\varepsilon \approx 0.4$ and ≈ 8 . The fit is to SC-ODT; the clock-relaxed inflow of § 6.3,
 1027 fitted identically, is carried alongside from figure 19 onward and the two are compared at




Figure 17. Change of the Amiet leading-edge-noise kernel when the frozen inflow spectrum is replaced by the computed distorted one, against reduced streamwise wavenumber, at total strains $e = 0.5$ – 2 (shade light to dark with increasing strain): (a) $Sk_t/\varepsilon \approx 0.4$, (b) ≈ 8 . Solid: spectral reallocation at matched upwash variance (shape only); dotted: total, including the measured variance change (-3.5 to -7.3 dB at the slow rate, where decay dominates the transit time; $+0.5$ to $+1.9$ dB at the rapid rate, where the upwash amplification survives). Geometry, aerofoil response and observer position cancel identically in this ratio.

figure 20. The line data pin the resolved-wavenumber structure; the transverse structure that (7.1) requires is supplied by the representation—this is § 5’s statement about line knowledge, operating in the application—and the representation’s rms residual, 13–24% over the strain range (the distance envelope of figure 12), is the corresponding uncertainty carried by the result. The two-dimensional spectrum follows by direct integration of the distorted components, $\Phi_{ww}(K_x, K_y; e) = A_0 \int \Phi_{22}^{\text{RDT}}(K_x, \kappa_2, K_y; e) d\kappa_2$, with the machinery verified against the wavevector-ensemble integrator of § 3.2 (upwash-variance growth reproduced to under 1%). The frozen baseline is the same inflow’s relaxed $e = 0$ state carried undistorted. The level and shape contributions are separated: the kernel scales linearly with the upwash variance, so dividing it out isolates the spectral reallocation, and the variance ratio itself is taken from the measured line spectra rather than from the fit. The chain is laid out step by step in the supplementary material (§ S4).

Figure 17 shows the result, and its sign and structure are the finding. Each panel plots, against reduced streamwise wavenumber $K_x/\kappa_e(0)$, the change in the noise kernel when the frozen inflow is replaced by the distorted one: the $\Delta\text{SPL} = 0$ line at the top is the frozen reference, so every curve lying below it marks a weaker source, and the curves darken with increasing total strain. The solid traces are the shape change at fixed upwash variance and the dotted traces add the measured variance term, so the offset between a solid curve and its dotted partner is that variance contribution—deepening the drop at the slow rate (panel a) and slightly lifting it at the rapid rate (panel b). At matched upwash variance the distorted spectrum is a substantially *weaker* leading-edge source than the frozen one at every resolved reduced frequency: ≈ -5 dB at $K_x \approx \kappa_e(0)$ for $e = 0.5$, deepening to -20 dB and beyond by $e = 2$, and increasingly severe towards low K_x . The mechanism is the orientation physics of § 5 acting on the noise kernel: (7.1) weights the spanwise-coherent upwash content $\Phi_{ww}(K_x, 0)$, which lives on wavevectors aligned with the line—exactly the orientations whose upwash is smallest by solenoidality and whose pre-images retreat into the vanishing low-wavenumber tail of the initial spectrum as the distortion proceeds.




Figure 18. The assembled absolute prediction at the reference configuration of § 4.4: far-field sound pressure level from the compact lift-dipole form of Amiet (1975) with the exact Sears response, for a flat plate of chord 0.39 m ($Re_c \approx 5 \times 10^5$) and span 0.8 m at $U_\infty = 20 \text{ m s}^{-1}$, $u'/U_\infty = 6\%$, $\mathcal{L} = 55 \text{ mm}$, observer 1 m overhead. Dashed: the frozen von Kármán inflow of standard practice (the $e = 0$ member of the fitted representation, which is exactly the isotropic von Kármán spectrum); grey dotted: the frozen Liepmann spectrum at the same $(u', \mathcal{L})$, an independent analytic cross-check of the physical-unit bridge. Solid (grey light to black with increasing strain, circles): the same chain with the SC-ODT-distorted inflow at total strains $e = 0.5, 1$ and 2 ($Sk_t/\varepsilon \approx 0.4$), level and shape both included.

The distortion, in short, drains the very content the aerofoil responds to. The total change adds the variance term: at the rapid rate the surviving upwash amplification offsets a small part of the reallocation; at the slow rate viscous decay over the longer transit time deepens it.

Figure 18 places the kernel result in the quantity practice actually reports: the absolute far-field spectrum of the reference configuration of § 4.4, assembled from the compact lift-dipole form of Amiet's theory with the exact Sears response (Sears 1941). The spectra enter in physical units through two matching conditions only—the integral scale fixes the length unit through the von Kármán relation $\kappa_e = 0.7468/\mathcal{L}$ and the upwash variance at $e = 0$ fixes the amplitude—so the frozen baseline *is* the standard Amiet–von Kármán prediction, not a recalibration of it: the $e = 0$ member of the representation is the isotropic von Kármán spectrum itself, and it reproduces the analytic one-dimensional form to within the quoted fit residual (ratio 0.79–1.14 across 100 Hz–1 kHz), with the Liepmann spectrum shown as an independent check of the unit bridge. The figure is read vertically, at fixed frequency: the two upper curves—dashed and dotted, all but coincident—are the frozen baselines, and their overlap is that unit-bridge check made visible; the three solid curves below are the same chain fed the distorted inflow at $e = 0.5, 1$ and 2 , so the downward offset of each from the dashed baseline is the level reduction the closure predicts at that strain, deepening as the strain increases. Against this baseline the computed distortion lowers the predicted level by 5 dB at $e = 0.5$, 9 dB at $e = 1$ and 14 dB at $e = 2$ across the resolved band ($f \lesssim 1.8 \text{ kHz}$ at this operating point), the spanwise-coherence drain of figure 17 now carried through geometry, aerofoil response and radiation to the observer. Both this figure and the kernel ratio of figure 17 carry the SC-ODT inflow as the representative computed spectrum—the clock-relaxed variant gives an indistinguishable far-field prediction here,




Figure 19. Absolute far-field predictions against published measurements, all four inflow spectra through the identical non-compact response (Amiet 1975; Roger & Moreau 2010): frozen von Kármán (black dashed), frozen Liepmann (grey dotted), the computed SC-ODT spectrum (grey, circles) and the computed clock-relaxed ODT spectrum (grey, triangles), each entered at the case’s effective free-distortion strain. (a,b) NACA0012 of Paterson & Amiet (1976) ($c = 0.23$ m, $u'_w/U \approx 4\text{--}5\%$, $\Lambda_f \approx 30$ mm, observer 2.25 m overhead; Gershfeld thickness admittance applied identically to all four chains); (c) ECL flat plate of Bampanis *et al.* (2022) ($c = 0.10$ m, 3 mm thick, $Tu = 4.5\%$, $\Lambda = 9$ mm, $\Theta = 90^\circ$, $R = 1.25$ m); (d) ISVR flat plate of Narayanan *et al.* (2015) ($c = 0.15$ m, $Tu = 2.5\%$, $\Lambda = 6$ mm, $R = 1.2$ m). Black measured symbols lie in each case’s stated comparison band (grey outside: background flags, facility contamination, thickness-breakdown frequency); spectra digitised from the published figures with per-figure calibration checks (± 0.5 dB). The effective strain is a geometric estimate (potential-flow stagnation-line accumulation truncated at the eddy scale, varied over $\Lambda/4\text{--}\Lambda$); the resulting spread in predicted level is at most ≈ 1 dB, comparable to the digitisation uncertainty, and is not shown.

for the reason quantified at figure 20—so the two computed inflows are carried separately only from the validation of figure 19 onward, where the comparison is the point.

The chain is then confronted with measurement. Figure 19 shows four published leading-edge-noise experiments with fully documented inflow turbulence—the NACA0012 of Paterson & Amiet (1976) at $40\text{--}90\text{ m s}^{-1}$ (higher speeds excluded only by their Mach number), and the flat plates of Bampanis *et al.* (2022) and Narayanan *et al.* (2015)—each with four predictions computed through the identical response: the frozen von Kármán




Figure 20. Computed against prescribed inflow at application strains, judged by the linear reference: absolute SPL at the reference configuration for the rapid family ($Sk_t/\varepsilon \approx 8$) at (a) $e = 1$ and (b) $e = 2$. Solid black: exact (inviscid) RDT of the model's relaxed initial state—the reference the strained-box DNS of § 4.6 validates at rapid rates—with level set by the kernel's exact upwash amplification (reproducing the moment-level value to 0.7%). Black dashed: the frozen von Kármán inflow of standard practice. Black with markers (circles: SC-ODT; triangles: clock-relaxed ODT): the two computed inflows (fitted representations, measured ensemble levels), each with a grey band giving its measured spectral rms distance from the RDT family mapped to decibels; the bands overlap sufficiently that the two inflows are indistinguishable at this level.

and Liepmann inflows of standard practice, and the model's computed spectra, SC-ODT and clock-relaxed ODT, evaluated at the effective free-distortion strain accumulated along the potential-flow stagnation streamline down to the eddy scale ($e_{\text{eff}} = 0.09\text{--}0.16$ for the NACA0012; below 0.05 for the thin plates). Three statements result. First, the assembled chain is quantitatively right where the experiment is cleanest: on the flat-plate cases every formulation sits within 1–2 dB rms of the measured spectra with the non-compactness interference lobes reproduced—the fidelity [Bampanis *et al.* \(2022\)](#) report for their own implementation—so the comparison discriminates inflow physics, not implementation. Second, the computed distortion behaves causally: for the thin plates the effective strain is negligible and both computed inflows correctly predict *no* correction to the frozen spectrum, while for the NACA0012 they lower the prediction towards the data at every speed (bias $+5.7 \rightarrow +5.3$ dB at 40 m s^{-1} , $+1.8 \rightarrow +1.4$ dB at 60, $+5.1 \rightarrow +4.7$ dB at 90), by the modest amount the geometric effective strain implies; the shift between predictions is exact, since the digitisation uncertainty (± 0.5 dB) is common to all four. Third, the SC-ODT and clock-relaxed ODT models are indistinguishable in the far field at these laboratory strains (differences below 0.05 dB), and this is stated as a limit of the test, not of the mechanism: laboratory rigs are designed to minimise inflow distortion, so no published far-field measurement samples the strains at which the relaxation mechanism acts. The residual NACA0012 overprediction is the finite-thickness effect that [Paterson & Amiet \(1976\)](#) themselves identified (breakdown near $U/ft_A \approx 1$) and that [Gershfeld \(2004\)](#) modelled, of a piece with the blocking mechanism excluded from the present scope; it is insensitive to the choice of inflow spectrum.

Where the mechanism does reach the prediction is shown in figure 20, and the claim is worded exactly as the evidence permits. Each panel shows absolute SPL against frequency




Figure 21. The cost of the frozen-inflow assumption against accumulated distortion: signed far-field prediction error, judged by the best available reference at each strain. Laboratory window (shaded): the eight measured cases of figure 19, bias over each case’s comparison band—flat plates filled, NACA0012 open (its common overprediction is the finite-thickness effect); horizontal bars span the effective-strain truncation range, and the three inflow chains are offset slightly in ϵ for legibility. Beyond the window: band-mean error over 100–1000 Hz against the exact-RDT reference of figure 20 at $e = 0.5, 1$ and 2 ; shaded bands give the $\pm 10 \log_{10}(1 + \text{rms})$ spectral-distance uncertainty where measured ($e = 1, 2$).

at one application strain; the black solid curve is the exact-RDT reference that the strained-box DNS validates at this rate, so the vertical distance of any other curve from it, read at fixed frequency, is that inflow’s error. The dashed frozen-von-Kármán curve sits above the reference—the over-prediction of standard practice—while the two computed inflows (circles, SC-ODT; triangles, clock-relaxed) fall on the reference within a shared grey band, close enough that they cannot be told apart at this level. At total strains of 1–2 the frozen-spectrum error against the linear reference grows from +1.3 to +3.6 dB, while both computed inflows track the reference within 1–2.5 dB—a residual that contains the physical relaxation at finite rapidity as well as the model deficiency measured in § 4.6.2—so an evolved inflow outperforms a prescribed one at the strains where the correction matters, whichever variant of the model evolves it. Between the two variants the radiated *level* does not discriminate: the clock acts below its size cap, where the noise kernel carries little weight, and its gain appears instead in the uncertainty the input carries into the prediction—the clock-relaxed ensembles lie closer to the exact-RDT-distorted family by a factor of two to four in rms at both rapidities, and, unlike SC-ODT’s, their distance does not grow with strain (supplementary material, § S3), which halves the shaded band of figure 20 and makes it strain-independent. A far-field discrimination of the mechanism requires a configuration whose acoustic response samples the sub-cap scales at accumulated strains of order unity—turbulence-ingesting rotors (Majumdar & Peake 1998), where the streamtube contraction supplies the strain and the blade response the fine-scale weighting—and is left as a falsifiable future target.

Figure 21 assembles the two regimes on a single axis and summarises the section. In the laboratory window the frozen and computed inflows are indistinguishable and both are right: the window in which Amiet’s theory has been validated is precisely the window in which its frozen-spectrum input costs nothing, which is why the practice has survived

unchallenged. What the methodology adds is the right-hand half of the figure: an inflow that continues to predict where the frozen assumption fails. Beyond the window the frozen error grows monotonically—+0.6, +1.4, +3.4 dB in the band mean at $e = 0.5, 1, 2$ —with no measurement there to catch it, while the computed inflows remain within 1.1–2.7 dB of the reference and carry their measured uncertainty with them.

The scope of this statement is set out as carefully as its content. It concerns the *free* distortion by the mean strain: the surface-blocking effect, which measurements and simulations show *enhancing* the near-edge upwash at low frequency (Hunt & Graham 1978; dos Santos *et al.* 2023; Piccolo *et al.* 2024), is a distinct and partially compensating mechanism outside the scope of the present line model; the calculation also replaces the strain history along individual streamlines by the canonical constant-rate distortion. What the calculation establishes is therefore not a noise prediction but the magnitude and sign of the term that frozen-spectrum practice omits: order 5–20 dB at the energy-containing reduced frequencies for total strains of 0.5–2—far too large to neglect, and of the opposite sign to the blocking correction with which it must ultimately be combined.

Figure 18 carries the kernel result through to an absolute far-field spectrum in physical units, the frozen baseline coinciding with the standard Amiet–von Kármán prediction; what remains for a predictive comparison against measurement is the blocking correction and the streamline strain history, for which the present model supplies the distorted-inflow input, with its uncertainty quantified by the envelope of § 4.6.2 and its fidelity improved by the mechanism of § 6 precisely in the wavenumbers with the greatest influence on the response.

8. Conclusions

We have presented a strain-coupled formulation of one-dimensional turbulence and, around it, a systematic account of what a single line of turbulence can and cannot say about rapid distortion. The formulation splits the distortion physics by structure: production enters as an exact continuous forcing, the rapid pressure–strain as an energy-conserving operator built from an explicit spectral closure, and the wavevector kinematics as a dilatation of the domain, while the eddy events of the standard model continue to carry the cascade and the slow return to isotropy. The construction preserves the model’s structural invariants, reproduces the constant-free rapid-distortion onset rate identically, and is verified computationally to the level of the closure it evolves (10^{-4} in the component fractions); at the moment level it reproduces the strained components of the Lee & Reynolds (1985) DNS quantitatively, with the one exception—the neutral-direction anisotropy—traced to the orientation information that every single-point closure discards, not to the line. The laboratory counterpart is the plane-strain experiment of Chen *et al.* (2006): the model reproduces the qualitative departure from rapid distortion that the experiment reports, the component centroids held far below the rigid-translation law and the upwash held back most (§ 4.5; supplementary material, § S1).

Three findings shape how such a model must be judged. First, the correct linear reference for line-projected spectra is broadband: exact rapid-distortion theory, projected onto a line, rescales the projected upwash spectrum by a factor of 2.4 across the resolved range at $e = 1$, so a rigid spectral translation—the model’s own kinematics—is not the linear prediction, and comparisons built on it conflate a structural approximation with nonlinear physics. Second, against a strained-box DNS with exact-RDT companions reduced to the same observables, the strain-coupled model carries the correct integrated anisotropy but allocates it uniformly over wavenumber, while linear theory and DNS concentrate it at

the energy-containing scales; the deficiency is localised in the wavenumber-uniform rapid operator and grows with strain rapidity, a measured artefact that accompanies the model into any application. Third, the single-line closure question admits an exact, falsifiable answer in place of a closure claim: line data bound the rapid term sharply, the bounds are valid up to $e \approx 0.25\text{--}0.5$ under plane strain and at all computed strains under axisymmetric distortion, and the axisymmetry hypothesis is measured rather than assumed—its plane-strain violation grows with accumulated strain and rapidity while the acoustically relevant upwash amplification remains robust to it.

The measured limits point to a specific missing ingredient that is thereupon identified and supplied. The strain-induced fine-scale anisotropy is delivered directly by large-eddy maps, which is why no acceptance-gating or map-reweighting scheme can control it; what does move it is supplementary relaxation applied at the scale of the imbalance—the model’s own energy-equalising kernel, applied hierarchically to eddy sub-scales. The intervention broadly reproduced the phenomenology: depth that reaches the measured scales is necessary and sufficient at moderate strain; dosage and delocalised application raise the early gains in proportion; and the fine-scale fade beyond $e \approx 3$, common to every scheme locked to the eddy clock, is removed by giving the relaxation its own clock—an independent event sequence of relaxation events whose size cap preserves the one-point physics. What remains is not timing but allocation: isotropisation is bounded at component equality and cannot reach the fine-scale sign reversal the DNS exhibits, so the remaining step is the wavenumber- and angle-dependent spectral kernel derived in § 5; exercising it in place of the wavenumber-uniform closure is the natural future continuation, now motivated by a measured deficit and pointed at a measured target.

For the application that motivated the work, the model delivers what the prediction chain otherwise lacks: a distorted, anisotropic, scale-resolved upwash spectrum, evolved rather than prescribed, at two orders of magnitude below the cost of a strained DNS ensemble—and § 7 carries it through Amiet’s theory to the radiated level. The assembled chain reproduces four published far-field experiments to 1–2 dB rms where the facilities are cleanest, with the computed correction causal—vanishing for thin plates, lowering the prediction towards the data for the NACA0012—and, against the linear reference at application strains, an evolved inflow outperforms the frozen one by a margin that grows with strain. Relative to the frozen-spectrum input of standard practice the computed distortion is a weaker leading-edge source at every resolved frequency—by 5–20 dB at the energy-containing scales for total strains 0.5–2, in the geometry-free kernel ratio and in the assembled absolute prediction at the reference configuration alike—because the distortion drains exactly the spanwise-coherent upwash content the aerofoil responds to. The omitted term is far too large to neglect and opposite in sign to the surface-blocking correction that a complete prediction must likewise include; predictions inherit a spectral-shape uncertainty of order the distance envelope at the operating point’s strain and rapidity, a falsifiable model prediction.

Appendix A. The forward map and the bound kernels

This appendix derives the three ingredients that § 5 uses: the two-scalar representation of the axisymmetric solenoidal spectrum with its realisability cone, the forward map from the two scalars to the line spectra (with the falsifiers T1–T2), and the closure kernels (5.7)–(5.9), together with the linear programs that turn these linear relations into the sharp bounds of § 5.4.

1230

A.1. The axisymmetric representation and realisability

The spectrum tensor of statistically homogeneous turbulence is Hermitian, positive semi-definite at each κ , and solenoidal, $\kappa_m \Phi_{mn}(\kappa) = 0$; it therefore has rank at most two, supported on the plane normal to κ . Upon invoking axisymmetry about $\mathbf{e}_2$, mirror symmetry, and zero helicity, here denoted Assumption A1, it is fixed by two scalar functions of $(\kappa_2, \kappa_\perp)$ (Batchelor 1946; Chandrasekhar 1950; Lindborg 1995):

$$\Phi_{mn}(\kappa) = a(\kappa_2, \kappa_\perp) P_{mn}(\kappa) + c(\kappa_2, \kappa_\perp) \xi_m \xi_n, \quad \xi = \mathbf{e}_2 - \mu \frac{\kappa}{\kappa}, \quad \mu = \frac{\kappa_2}{\kappa}, \quad (\text{A } 1)$$

with $P_{mn} = \delta_{mn} - \kappa_m \kappa_n / \kappa^2$ the solenoidal projector, so that $\kappa_m \xi_m = 0$ and $|\xi|^2 = 1 - \mu^2 = 1 - x$ with $x = \kappa_2^2 / \kappa^2$: solenoidality is exact by construction, which is how the third would-be scalar of a general axisymmetric tensor is eliminated. For a unit vector $\mathbf{v}$ normal to κ , $\Phi_{mn} v_m v_n = a + c(\xi \cdot \mathbf{v})^2$, so the two eigenvalues of the polarisation matrix on that plane are a (attained for $\mathbf{v} \perp \xi$) and $a + (1 - x)c$ (for $\mathbf{v} \parallel \xi$). Positive semi-definiteness is therefore the pointwise *realisability cone*

$$a(\kappa_2, \kappa_\perp) \geq 0, \quad a(\kappa_2, \kappa_\perp) + (1 - x)c(\kappa_2, \kappa_\perp) \geq 0, \quad (\text{A } 2)$$

1231 quoted in § 5.4. Isotropy is the special case $c = 0$ with a a function of κ alone.

1232

1233

A.2. The forward map to the line spectra

Writing $\kappa = (\kappa_\perp \cos \theta, \kappa_2, \kappa_\perp \sin \theta)$, the azimuthal moments at fixed $(\kappa_2, \kappa_\perp)$ are $\langle \kappa_1^2 \rangle_\theta = \langle \kappa_3^2 \rangle_\theta = \frac{1}{2} \kappa_\perp^2$, $\langle \kappa_1^4 \rangle_\theta = \langle \kappa_3^4 \rangle_\theta = \frac{3}{8} \kappa_\perp^4$, $\langle \kappa_1^2 \kappa_3^2 \rangle_\theta = \frac{1}{8} \kappa_\perp^4$, and every odd moment vanishes. Applied to (A 1) the diagonal weights average to

$$\langle P_{11} \rangle_\theta = \langle P_{33} \rangle_\theta = \frac{1}{2}(1 + x), \quad P_{22} = 1 - x, \quad \langle \xi_1^2 \rangle_\theta = \langle \xi_3^2 \rangle_\theta = \frac{1}{2}x(1 - x), \quad \xi_2^2 = (1 - x)^2, \quad (\text{A } 3)$$

and every mixed weight to zero. Substituting into the definition (5.3) of the line spectra, $\phi_{mn}(\kappa_2) = \int_0^\infty \langle \Phi_{mn} \rangle_\theta 2\pi \kappa_\perp d\kappa_\perp$, gives the **forward map** from the two scalars to the data a line carries:

$$\phi_{11}(\kappa_2) = \phi_{33}(\kappa_2) = 2\pi \int_0^\infty \left[a \frac{1+x}{2} + c \frac{x(1-x)}{2} \right] \kappa_\perp d\kappa_\perp, \quad (\text{A } 4)$$

$$\phi_{22}(\kappa_2) = 2\pi \int_0^\infty \left[a(1 - x) + c(1 - x)^2 \right] \kappa_\perp d\kappa_\perp, \quad (\text{A } 5)$$

$$\phi_{mn}(\kappa_2) = 0 \quad (m \neq n). \quad (\text{A } 6)$$

1234 Equations (A 4)–(A 6) contain the two free falsifiers quoted in § 5.4: under A1 the two transverse line spectra
1235 coincide at every resolved wavenumber (T1), and the cross-spectra vanish (T2); both are properties the line itself
1236 reports, scale by scale, with no reference to unresolved data. The ϕ_{22} bracket is the same combination of a and
1237 c that appears in the kernel g_2 below—the conventions mesh, as they must, because ϕ_{22} and the 22-component
1238 of the rapid term weight the spectrum through the same projection geometry.

1239 They also explain why the isotropic intuition about line data fails under A1. For isotropic turbulence ($c = 0$,
1240 $a = a(\kappa)$) the map is over-determined: one unknown function of one variable against two measured functions,
1241 with ϕ_{22} determining a by the classical Abel-type inversion and the transverse spectrum slaved to it by the
1242 derivative relation. Under A1 the bookkeeping reverses: *two unknown functions of two variables* against two
1243 measured functions of one variable. At each fixed κ_2 the data supply exactly two numbers—the two weighted
1244 $\kappa_\perp$ -integrals (A 4)–(A 5)—about two unknown functions of $\kappa_\perp$; any perturbation $(\delta a, \delta c)(\kappa_\perp)$ annihilating both
1245 integrals while respecting (A 2) is invisible to the line. This infinite-dimensional null space is the quantitative
1246 content of the obstruction of § 5.2 after the axisymmetric collapse.

1247

1248

A.3. Derivation of the closure kernels

For the diagonal components the exact result (5.2) reads

$$\Pi_{nn}^{(r)} = 2 A_{kl} \int_{\mathbb{R}^3} \frac{\kappa_n \kappa_k}{\kappa^2} \Phi_{nl}(\kappa) d\kappa \quad (\text{no sum on } n). \quad (\text{A } 7)$$

Substituting (A 1), carrying out the azimuthal average with the moments above, and specialising to $A = \text{diag}(\frac{1}{2}, -\frac{1}{2}, 0)$ reduces each component to a $2\pi \kappa_\perp$ -weighted integral over the half-plane of a polynomial combination in x , linear in (a, c) —the kernels (5.7)–(5.9) of the main text. The upwash component illustrates

the mechanics: only $k = l = 1$ and $k = l = 2$ contribute, with azimuthal averages $\langle \kappa_1 \Phi_{21} \rangle_\theta = -\frac{1}{2} \kappa_\perp^2 \kappa_2 [a + (1-x)c]/\kappa^2$ and $\langle \Phi_{22} \rangle_\theta = (1-x)[a + (1-x)c]$, so that both enter through the same bracket and

$$\left\langle \text{integrand of (A 7)} \right\rangle_\theta \Big|_{n=2} = 2 \left[-\frac{1}{2} \cdot \frac{x(1-x)}{2} - \frac{1}{2} x(1-x) \right] [a + (1-x)c] = -\frac{3}{2} x(1-x)[a + (1-x)c], \quad (\text{A 8})$$

which is $g_2(x)$. The transverse components involve in addition the fourth moments of (A 3)’s underlying table and yield g_1 and g_3 ; the full reduction has been verified symbolically against (A 7). Three identities check the algebra: $g_1 + g_2 + g_3 = 0$ at every x for both scalars (pointwise energy conservation); each kernel vanishes at $x = 1$, where the wavevector aligns with the line and the projection degenerates; and averaging over the isotropic angular measure $d\mu$ on $[0, 1]$ gives $\langle g_1 \rangle : \langle g_2 \rangle : \langle g_3 \rangle = +\frac{1}{5} : -\frac{1}{5} : 0$, reproducing $\Pi_{ij}^{(r)} = \frac{4}{5} k_i S_{ij}$ (Crow 1968) as in § 5.7.

A.4. The bounds as linear programs

The target (5.6), the data constraints (A 4)–(A 5), and the cone (A 2) are all linear in (a, c) , and everything decomposes over the resolved wavenumber: writing $\Pi_{nn}^{(r)} = \int_0^\infty \pi_n(\kappa_2) d\kappa_2$ with $\pi_n = 2\pi \int_0^\infty g_n(x) [a, c] \kappa_\perp d\kappa_\perp$, the extremisation over all spectra consistent with the line data is a family of *independent* infinite-dimensional linear programs, one per κ_2 :

$$\pi_n^\pm(\kappa_2) = \max / \min_{(a,c)(\kappa_\perp)} 2\pi \int_0^\infty g_n(x) [a, c] \kappa_\perp d\kappa_\perp \quad \text{s.t.} \quad (\text{A 4})\text{--}(\text{A 5}) \text{ hold, (A 2) pointwise,} \quad (\text{A 9})$$

whence $\Pi_{nn}^\pm = \int \pi_n^\pm d\kappa_2$ and the indeterminacy $\mathcal{I}_n = (\Pi_{nn}^+ - \Pi_{nn}^-)/2k_i S$ of § 5.4, normalised by the natural rapid-term scale so that $\mathcal{I}_n$ is comparable across strains (the isotropic value is $\Pi_{11}^{(r)} = \frac{2}{5} k_i$ per unit $S_{11} = \frac{1}{2}$). Discretised on a logarithmic $\kappa_\perp$ grid of order 10^2 points, (A 9) is a linear program with two equality constraints and the cone, solved per κ_2 in milliseconds; the bounds are global and sharp by LP duality—no fitting and no local minima. The side conditions of § 5.4 enter as further linear constraints: the log-Lipschitz slope cap $|d \ln(a, c)/d \ln \kappa_\perp| \leq \lambda$ and the polarisation cap $|c| \leq \gamma a$. The LP extremals concentrate on few-point supports in $\kappa_\perp$ (bang–bang spectra), which is why the raw bounds are conservative and why smoothness alone barely narrows them: a log-Lipschitz spectrum can still vary like $\kappa_\perp^{\pm \lambda}$, so the angular reweighting that drives the indeterminacy survives smoothing, while the polarisation cap attacks it directly. The dual variables of (A 9) certify, per resolved wavenumber, which linear combinations of the measured line spectra control each π_n ; finite-window bias on ϕ_{mn} at small κ_2 propagates linearly and can be carried as interval data in place of the equalities with no change of machinery.

Appendix B. Numerical parameters, initial conditions and statistics

B.1. Model configurations and intervention labels

The reference configurations and every intervention variant of the measured-limits campaign (§ 6; supplementary material § S2) carry a *category.variant* label, used with that label throughout. The reference configurations are C0.1, standard ODT (the strain-off limit $A = 0$), and C0.2, SC-ODT (the strain-coupled model the interventions modify); the intervention categories C1–C5 are listed in table 4. The three headline models of the main text are C0.1 (standard ODT), C0.2 (SC-ODT) and C5.2 (clock-relaxed ODT), and the five-way comparison of figure 16 is C0.1, C0.2 and C5.2 with exact RDT and DNS. The size cap of the clock-relaxed variant (C5.2) is $\ell^* = 0.05$ on the unit domain, the scale separating the two transmission-diagnostic bands of § 6.

B.2. ODT configurations

All line ensembles share the solver settings of the standard model (Stephens & Lignell 2021): eddy-rate constant $C = 7$, viscous-penalty parameter $Z = 450$, kernel parameter $\alpha = 2/3$, adaptive mesh with $dx \in [2 \times 10^{-4}, 2 \times 10^{-2}]$ on a unit domain (4000 initial cells), explicit advancement with diffusive CFL 0.5 and a strain-resolved bound of 10^{-2} on the total-strain increment per step (§ 3.3). The initial condition is the resolved spectral field of § 3.3: a band-limited isotropic superposition of 64 Fourier modes peaked at eight wavelengths per domain, which places no energy at the grid scale. Strained production ensembles use an unstrained precursor:

label	configuration / intervention (knob)
C0.1	standard ODT (strain-off limit, $A = 0$)
C0.2	SC-ODT (strain-coupled; the configuration the interventions modify)
<i>C1 acceptance gating</i>	
C1.1	gate on post-event anisotropy (45–98% rejection)
C1.2	gate on sub-band eddies only (42%, 97% rejection)
<i>C2 map-image geometry</i>	
C2.1	unequal triplet-map images (middle image of relative size $\frac{2}{3}$)
<i>C3 sub-scale relaxation kernels (depth)</i>	
C3.1	one level ($\ell/3$)
C3.2	two levels ($\ell/9$)
C3.3	three levels ($\ell/27$)
C3.4	recursion depth tied to eddy size
C3.5	two-level pass iterated $3\times$ per event
<i>C4 concurrent relaxation</i>	
C4.1	$N = 1$ randomly placed interval per event
C4.2	$N = 3$ randomly placed intervals per event
C4.3	concurrent + depth: adaptive hierarchy in each of $N = 3$ intervals (stacked)
<i>C5 independent relaxation clock</i>	
C5.1	uncapped (rate matched to the eddy rate)
C5.2	size-capped $\ell \leq \ell^*$, adaptive depth inside — the clock-relaxed ODT
C5.3	size-capped clock at $4\times$ the eddy rate

Table 4. Category.variant labels for the reference configurations and the intervention categories. Measured fine-scale effects and remarks are in supplementary table S1; the four-panel diagnostic is figure S3.

the line evolves as standard ODT until the eddy population is statistically active and the initialisation transient has relaxed, and the strain is switched on only then (§ 4.6). The molecular viscosity is set per study: 10^{-4} for the transmission diagnostic of § 6 (a scale diagnostic, resolved with wide margin), 10^{-5} for the strain-on/off ensembles of § 4.3, and 3×10^{-6} on a 16000-cell grid at the operating point of § 4.4. The strained line spectrum ensembles carried into the leading-edge-noise chain of Part III and the relaxation-clock spectra (supplementary material § S3) are 1024 realisations per case at each rapidity $Sk_t/\varepsilon \approx 0.4$ and ≈ 8 . The resolved scale separation of the $\nu = 10^{-5}$ ensembles, used for the CMK comparison of supplementary § S1, is $\mathcal{L}/\eta \approx 250$.

B.3. Strained-box DNS and exact-RDT companions

The benchmark of § 4.6 is a pseudo-spectral 128^3 simulation in a Rogallo deforming frame: a decaying isotropic precursor, then plane strain $A = S \text{diag}(\frac{1}{2}, -\frac{1}{2}, 0)$ at $Sk_t/\varepsilon \in \{0.8, 16\}$ at onset, carried to $e = 1$ over four independent realisations, with $\kappa_{\max}\eta \approx 0.76$ at $e = 1$ —a pilot resolution, so small-scale statements are qualitative. Every realisation has an exact-RDT companion: the same initial field evolved under the closed-form Cauchy solution of the linearised equations and projected onto lines identically. The initial Fourier fields are whitened by the *symmetric* inverse square root $C^{-1/2}$ of the sample covariance, so the initial anisotropy vanishes per realisation. This detail is load-bearing: an ordered-Cholesky whitening used initially left a bias $b_{33} = +0.0148 \pm 0.0046$, detected by the nominal-turbulence null test (strain machinery active, $A = 0$) and traced by a reversal experiment to the ordering; with the symmetric form all $|b_{ii}| \leq 0.005$ and the on-line

falsifier T1 exceeds its confidence band at the chance rate. The estimator pipeline of § 5.4 passes the same null battery on JHTDB isotropic turbulence (1024³, 128 lines): T1 exceedances at the chance rate, cross-spectral coherences $\sim 10^{-2}$, the longitudinal–transverse derivative relation to 4%, and the LP bounds bracketing the exact isotropic value at every constraint level.

B.4. Statistics policy

Spectral band quantities of the strained line are strongly intermittent across realisations—single realisations carry 10^2 – 10^3 times the median band energy—so ensemble means of band ratios do not converge at any practical sample size. All ODT spectral statistics are therefore medians over realisations with 95% bootstrap confidence intervals of the median, and every intervention in § 6 is assessed by the median of *paired* same-seed contrasts, quoted significant only when the bootstrap interval excludes zero. Ensemble sizes: 1024 realisations for the transmission diagnostic and the distance envelope, 128–200 for the spectrum ensembles of § 4, and four seeds for the DNS and its companions, whose derived quantities carry jackknife errors over realisations.

REFERENCES

- AMIET, R. K. 1975 Acoustic radiation from an airfoil in a turbulent stream. *Journal of Sound and Vibration* **41** (4), 407–420.
- ASHURST, WILLIAM T. & KERSTEIN, ALAN R. 2005 One-dimensional turbulence: variable-density formulation and application to mixing layers. *Physics of Fluids* **17** (2), 025107.
- BAMPANIS, G., ROGER, M. & MOREAU, S. 2022 On a three-dimensional investigation of airfoil turbulence-impingement noise and its reduction by leading-edge tubercles. *Journal of Sound and Vibration* **520**, 116635.
- BATCHELOR, G. K. 1946 The theory of axisymmetric turbulence. *Proceedings of the Royal Society of London. Series A* **186** (1007).
- BATCHELOR, G. K. & PROUDMAN, IAN 1954 The effect of rapid distortion of a fluid in turbulent motion. *The Quarterly Journal of Mechanics and Applied Mathematics* **7** (1), 83–103.
- CHANDRASEKHAR, S. 1950 The theory of axisymmetric turbulence. *Philosophical Transactions of the Royal Society of London. Series A* **242** (855).
- CHEN, JUN, MENEVEAU, CHARLES & KATZ, JOSEPH 2006 Scale interactions of turbulence subjected to a straining–relaxation–destraining cycle. *Journal of Fluid Mechanics* **562**, 123–150.
- CLAY, M. P. & YEUNG, P. K. 2016 A numerical study of turbulence under temporally evolving axisymmetric contraction and subsequent relaxation. *Journal of Fluid Mechanics* **805**, 460–493.
- CROW, S. C. 1968 Viscoelastic properties of fine-grained incompressible turbulence. *Journal of Fluid Mechanics* **33** (1), 1–20.
- GERSHFELD, J. 2004 Leading edge noise from thick foils in turbulent flows. *Journal of the Acoustical Society of America* **116** (3), 1416–1426.
- GOLDSTEIN, M. E. & ATASSI, H. 1976 A complete second-order theory for the unsteady flow about an airfoil due to a periodic gust. *Journal of Fluid Mechanics* **74** (4), 741–765.
- GUALTIERI, P. & MENEVEAU, C. 2010 Direct numerical simulations of turbulence subjected to a straining and destraining cycle. *Physics of Fluids* **22** (6), 065104.
- HUNT, J. C. R. 1973 A theory of turbulent flow round two-dimensional bluff bodies. *Journal of Fluid Mechanics* **61** (4), 625–706.
- HUNT, J. C. R. & CARRUTHERS, D. J. 1990 Rapid distortion theory and the ‘problems’ of turbulence. *Journal of Fluid Mechanics* **212**, 497–532.
- HUNT, J. C. R. & GRAHAM, J. M. R. 1978 Free-stream turbulence near plane boundaries. *Journal of Fluid Mechanics* **84** (2), 209–235.
- VON KÁRMÁN, THEODORE 1948 Progress in the statistical theory of turbulence. *Proceedings of the National Academy of Sciences* **34** (11), 530–539.
- KASSINOS, S. C. & REYNOLDS, W. C. 1998 A structure-based model with stropholysis effects. In *Annual Research Briefs 1998*, pp. 197–209. Center for Turbulence Research, Stanford University.
- KASSINOS, S. C., REYNOLDS, W. C. & ROGERS, M. M. 2001 One-point turbulence structure tensors. *Journal of Fluid Mechanics* **428**, 213–248.
- KERSTEIN, ALAN R. 1991 Linear-eddy modelling of turbulent transport. Part 6. Microstructure of diffusive scalar mixing fields. *Journal of Fluid Mechanics* **231**, 361–394.
- KERSTEIN, ALAN R. 1999 One-dimensional turbulence: model formulation and application to homogeneous turbulence, shear flows, and buoyant stratified flows. *Journal of Fluid Mechanics* **392**, 277–334.
- KERSTEIN, ALAN R., ASHURST, WILLIAM T., WUNSCH, SCOTT & NILSEN, VEBJORN 2001 One-dimensional turbulence: vector formulation and application to free shear flows. *Journal of Fluid Mechanics* **447**, 85–109.

- 1363 KIM, JAE WOOK & HAERI, SINA 2015 An advanced synthetic eddy method for the computation of aerofoil–
1364 turbulence interaction noise. *Journal of Computational Physics* **287**, 1–17.
- 1365 KOLMOGOROV, A. N. 1991 The local structure of turbulence in incompressible viscous fluid for very large
1366 Reynolds numbers. *Proceedings of the Royal Society of London. Series A* **434** (1890), 9–13.
- 1367 LAUNDER, B. E., REECE, G. J. & RODI, W. 1975 Progress in the development of a reynolds-stress turbulence
1368 closure. *Journal of Fluid Mechanics* **68** (3), 537–566.
- 1369 LEE, MOON J., PIOMELLI, UGO & REYNOLDS, WILLIAM C. 1986 Useful formulas in the rapid distortion theory
1370 of homogeneous turbulence. *Physics of Fluids* **29** (10), 3471–3474.
- 1371 LEE, MOON J. & REYNOLDS, WILLIAM C. 1985 Numerical experiments on the structure of homogeneous
1372 turbulence. *Tech. Rep.* TF-24. Thermosciences Division, Department of Mechanical Engineering, Stanford
1373 University, Stanford, California, prepared under NASA Grant NCC-2-15.
- 1374 LIGNELL, DAVID O., KERSTEIN, ALAN R., SUN, GUANGYUAN & MONSON, ELIZABETH I. 2013 Mesh adaption for
1375 efficient multiscale implementation of one-dimensional turbulence. *Theoretical and Computational Fluid*
1376 *Dynamics* **27** (3), 273–295.
- 1377 LIGNELL, DAVID O., LANSINGER, VICTORIA B., MEDINA MÉNDEZ, JUAN A., KLEIN, MARTEN, KERSTEIN, ALAN R.,
1378 SCHMIDT, HEIKO, FISTLER, MARCO & OEVERMANN, MICHAEL 2018 One-dimensional turbulence modeling
1379 for cylindrical and spherical flows: model formulation and application. *Theoretical and Computational Fluid*
1380 *Dynamics* **32** (4), 495–520.
- 1381 LINDBORG, E. 1995 Kinematics of homogeneous axisymmetric turbulence. *Journal of Fluid Mechanics* **302**,
1382 179–201.
- 1383 MAJUMDAR, S. J. & PEAKE, N. 1998 Noise generation by the interaction between ingested turbulence and a
1384 rotating fan. *Journal of Fluid Mechanics* **359**, 181–216.
- 1385 MEDINA MÉNDEZ, JUAN A., SHARMA, SPARSH, SCHMIDT, HEIKO & KLEIN, MARTEN 2023 Toward the use of a
1386 reduced-order and stochastic turbulence model for assessment of far-field sound radiation: Low Mach number
1387 jet flows. *Proceedings in Applied Mathematics and Mechanics* **23** (3), e202300186.
- 1388 MILLS, R. R. & CORRSIN, S. 1959 Effect of contraction on turbulence and temperature fluctuations generated
1389 by a warm grid. *NASA Memo* (5-5-59W).
- 1390 NARAYANAN, S., CHAITANYA, P., HAERI, S., JOSEPH, P., KIM, J. W. & POLACSEK, C. 2015 Airfoil noise reductions
1391 through leading edge serrations. *Physics of Fluids* **27** (2), 025109.
- 1392 PATERSON, R. W. & AMIET, R. K. 1976 Acoustic radiation and surface pressure characteristics of an airfoil due
1393 to incident turbulence. *Tech. Rep.* CR-2733. NASA.
- 1394 PATERSON, R. W. & AMIET, R. K. 1977 Noise and surface pressure response of an airfoil to incident turbulence.
1395 *Journal of Aircraft* **14** (8), 729–736.
- 1396 PICCOLO, A., ZAMPONI, R., AVALLONE, F. & RAGNI, D. 2024 Turbulence distortion and leading-edge noise.
1397 *Physics of Fluids* **36** (12), 125183.
- 1398 POPE, STEPHEN B. 2000 *Turbulent Flows*. Cambridge: Cambridge University Press.
- 1399 REYNOLDS, W. C. & KASSINOS, S. C. 1995 One-point modelling of rapidly deformed homogeneous turbulence.
1400 *Proceedings of the Royal Society of London A* **451** (1941), 87–104.
- 1401 RIBEIRO, A. F. P., DOS SANTOS, F. L., VENNER, C. H. & DE SANTANA, L. D. 2023 Numerical study of inflow
1402 turbulence distortion and noise for airfoils. *Physics of Fluids* **35** (11), 115112.
- 1403 ROGER, MICHEL & MOREAU, STÉPHANE 2010 Extensions and limitations of analytical airfoil broadband noise
1404 models. *International Journal of Aeroacoustics* **9** (3), 273–305.
- 1405 SAGAUT, P. & CAMBON, C. 2018 *Homogeneous Turbulence Dynamics*, 2nd edn. Cham: Springer.
- 1406 DE SANTANA, LEANDRO D., CHRISTOPHE, JULIEN, SCHRAM, CHRISTOPHE F. & DESMET, WIM 2016 A rapid
1407 distortion theory modified turbulence spectra for semi-analytical airfoil noise prediction. *Journal of Sound*
1408 *and Vibration* **383**, 349–363.
- 1409 DOS SANTOS, F. L., BOTERO-BOLIVAR, L., VENNER, C. H. & DE SANTANA, L. D. 2023 Inflow turbulence distortion
1410 for airfoil leading-edge noise prediction for large turbulence length scales for zero-mean loading. *The Journal*
1411 *of the Acoustical Society of America* **153** (3), 1811–1822.
- 1412 SCHMIDT, RODNEY C., KERSTEIN, ALAN R. & McDERMOTT, RANDALL 2010 ODTLES: A multi-scale model for
1413 3D turbulent flow based on one-dimensional turbulence modeling. *Computer Methods in Applied Mechanics*
1414 *and Engineering* **199** (9–12), 865–880.
- 1415 SEARS, WILLIAM R. 1941 Some aspects of non-stationary airfoil theory and its practical application. *Journal of*
1416 *the Aeronautical Sciences* **8** (3), 104–108.
- 1417 SHARMA, SPARSH, KLEIN, MARTEN & SCHMIDT, HEIKO 2022 Features of far-downstream asymptotic velocity
1418 fluctuations in a round jet: A one-dimensional turbulence study. *Physics of Fluids* **34** (8), 085131.
- 1419 STEPHENS, VICTORIA B. & LIGNELL, DAVID O. 2021 One-dimensional turbulence (ODT): Computationally
1420 efficient modeling and simulation of turbulent flows. *SoftwareX* **13**, 100641.

- 1421 TOWNSEND, A. A. 1976 *The Structure of Turbulent Shear Flow*, 2nd edn. Cambridge: Cambridge University
1422 Press.
- 1423 TUCKER, H. J. & REYNOLDS, A. J. 1968 The distortion of turbulence by irrotational plane strain. *Journal of*
1424 *Fluid Mechanics* **32** (4), 657–673.
- 1425 WUNSCH, SCOTT & KERSTEIN, ALAN R. 2001 A model for layer formation in stably stratified turbulence. *Physics*
1426 *of Fluids* **13** (3), 702–712.
- 1427 ZUSI, C. J. & PEROT, J. B. 2013 Simulation and modeling of turbulence subjected to a period of uniform plane
1428 strain. *Physics of Fluids* **25** (11), 110819.
- 1429 ZUSI, C. J. & PEROT, J. B. 2014 Simulation and modeling of turbulence subjected to a period of axisymmetric
1430 contraction or expansion. *Physics of Fluids* **26** (11), 115103.
- 1431